

Welcome to the Course

CompTIA SecAI+ | Section 1: Getting Started

Why AI Security Is the Most Important Skill in Tech Right Now

Artificial intelligence has moved from research laboratories into critical infrastructure. Healthcare systems now use AI to diagnose cancer from imaging data. Banks rely on ML models to approve loans and detect fraud in milliseconds. Autonomous vehicles make life-or-death decisions using neural networks. Power grids, air traffic control systems, and logistics networks are increasingly AI-driven.

With this expansion comes a new and rapidly growing attack surface. When an AI system is compromised, the consequences are fundamentally different from traditional software breaches. A poisoned fraud detection model can allow millions in losses before anyone notices. A manipulated medical diagnosis AI can harm patients. A jailbroken AI assistant in a corporate environment can leak confidential data to any user who knows the right prompt. These are not theoretical scenarios — they are documented, real incidents happening today.

AI Security Incident Type	Real-World Example	Business Impact	Why Traditional Security Missed It
Data poisoning	Microsoft Tay chatbot (2016): users poisoned the training feedback loop, causing the model to generate offensive outputs within 24 hours	Immediate product shutdown; reputational damage	No signature for adversarial training input; appeared as normal user interaction
Model evasion	Stop sign adversarial patches fooling autonomous vehicle vision systems (research, 2017+)	Vehicle misclassification of traffic signs at speed	Input looked like a valid stop sign to camera sensors and human eyes
Deepfake fraud	Arup Engineering CEO deepfake video call resulting in HK\$200M transfer (2024)	Direct financial loss; identity verification failure	No malware involved; attack exploited trust in visual and audio confirmation
Prompt injection	Samsung employees inadvertently leaked proprietary source code into ChatGPT training pipeline (2023)	Confidential IP exposure; regulatory risk	No network intrusion; data left via user-initiated legitimate API calls
AI supply chain attack	ShadowRay vulnerability: attackers exploited Ray AI framework to compromise ML workloads (2024)	Unauthorised access to AI training infrastructure	Exploit targeted AI-specific framework, not traditional application layer

The CompTIA SecAI+ Exam Domains: A Study Roadmap

The SecAI+ certification is structured around four exam domains that map directly to the competencies required in real AI security roles. Understanding what each domain covers and how much it weighs in the exam allows you to allocate your study time strategically.

Domain	Title	Exam Weight	Core Competencies Covered
Domain 1	AI Fundamentals	14%	ML types and algorithms, neural network architectures, model training and inference, data pipelines, statistical learning concepts
Domain 2	AI Security Operations	29%	Monitoring AI systems in production, detecting anomalous model behaviour, securing AI infrastructure, incident response for AI-specific incidents, SIEM and SOAR integration
Domain 3	AI Attack Vectors and Defense	38%	Adversarial attacks (evasion, poisoning, extraction, inversion), prompt injection and jailbreaking, deepfake detection, model hardening, red-teaming AI systems
Domain 4	Governance, Risk, and Compliance	19%	AI-specific risk frameworks including NIST AI RMF, EU AI Act and global regulations, ethical AI principles, model auditing, bias and fairness assessment, third-party AI vendor risk

Domain 3 (AI Attack Vectors and Defense) carries the highest weighting at 38%. This reflects the industry reality that understanding offensive AI techniques is essential for building effective defences. Students who come from a traditional cybersecurity background often find Domain 1 (AI Fundamentals) the most challenging — invest proportionally more time there if ML mathematics is new to you.

Who This Certification Is For: Career Pathways

SecAI+ sits at the intersection of two rapidly growing fields. The credential is designed for professionals who work at that intersection and need to understand both AI and security well enough to secure AI systems effectively.

Starting Background	Current Role	SecAI+ Opens These Paths	Typical Exam Prep Timeline
Cybersecurity	SOC Analyst, Penetration Tester, Security Engineer	AI Security Analyst, AI Red Team Lead, AI Security Architect	2 to 4 months (strong security base; focus on AI fundamentals)
IT and Systems	Systems Administrator, Network Engineer, Cloud Engineer	AI Infrastructure Security Engineer, AI Governance Specialist	3 to 5 months (learn both AI and security domains)
AI and Data Science	ML Engineer, Data Scientist, Data Analyst	AI Risk Manager, MLSecOps Engineer, AI Compliance Lead	2 to 3 months (strong AI base; focus on security fundamentals)
Risk and Compliance	GRC Analyst, Compliance Officer, Risk Manager	AI Governance Engineer, AI Ethics Officer, AI Regulatory Compliance	3 to 5 months (learn technical AI and security concepts)

		Lead	
Career Changer or Student	Other field or entry-level	AI Security Analyst (entry level), AI Security Intern, SOC Analyst with AI focus	5 to 8 months (build both domains from foundation)

Study Strategy: How to Learn This Material Effectively

AI security is a synthesis field. You cannot learn it through passive watching. The following strategies are grounded in cognitive science research on technical skill acquisition and have been validated by certification communities across comparable credentials.

- **Active recall over re-reading:** After each lecture, close your notes and attempt to write down everything you remember. The act of retrieving information from memory strengthens retention far more than re-reading notes. Use flash cards for definitions and acronyms.
- **Spaced repetition:** Review material at increasing intervals. After Day 1, review on Day 3, then Day 7, then Day 14. Tools like Anki implement this automatically. This technique can reduce the study hours needed to achieve the same retention level by 40 to 60%.
- **Concept mapping for AI security:** Many SecAI+ concepts are interconnected. Draw maps showing relationships between topics such as adversarial examples, ML architecture, training data quality, and governance. Spatial organisation helps pattern recognition during exams.
- **Apply concepts to current news:** AI security incidents are reported weekly. When you study prompt injection, search for the latest documented cases. Connecting abstract concepts to real events significantly accelerates understanding.
- **Teach-back method:** After mastering a topic, explain it aloud as if teaching a colleague with no background. Gaps in your explanation reveal gaps in your understanding before the exam reveals them instead.
- **Practice questions from day one:** Do not wait until content is complete before attempting practice questions. Attempting questions on recently covered material — and reviewing explanations of wrong answers in detail — is one of the most efficient study methods known.

The AI Security Skills Gap: Why This Credential Has Outsised Value Now

Certifications derive their market value from the ratio of employer demand to credential supply. When demand far exceeds supply, a certification commands premium salaries and provides significant competitive advantage. In 2026, AI security represents one of the most extreme examples of this imbalance in the technology sector.

According to ISC2 research, 89% of organisations report lacking sufficient AI security expertise — yet 70% of enterprises have deployed AI in production systems (McKinsey 2026). LinkedIn data shows AI security roles growing at 45% annually, significantly outpacing overall cybersecurity growth. The World Economic Forum identified AI security as one of the top 10 critical skills shortfalls of 2025 to 2030.

Market Reality: The SecAI+ is a new certification in a new field. Being an early adopter of a credential that becomes an industry standard is disproportionately valuable. The window for early-adopter advantage in AI security is open right now.

The Four Learning Habits That Predict Certification Success

Across technical certification communities, four behavioural patterns consistently distinguish candidates who pass on their first attempt from those who require multiple attempts.

Habit	What It Looks Like	Why It Works	How to Implement
Consistency over intensity	30 to 60 minutes daily rather than 8-hour weekend cramming sessions	Sleep consolidates memory; daily repetition builds neural pathways that enable rapid recall under exam pressure	Schedule a non-negotiable study slot in your calendar; treat it like a work meeting
Note-taking in your own words	Summarising concepts in your own language rather than copying slide text verbatim	Paraphrasing requires active processing; passive copying does not encode the concept into long-term memory	Use the Cornell method: main notes on the right, key questions on the left, summary at the bottom
Deliberate review of wrong answers	Spending more time on incorrectly answered practice questions than correct ones	Your wrong answers map precisely to your knowledge gaps; correct answers confirm what you already know	After each practice session, categorise every wrong answer by domain and create a personal weakness list
Application to real scenarios	Finding examples of studied concepts in security news, job descriptions, and tool documentation	Concrete examples create retrieval hooks that abstract definitions lack	Follow AI security-focused accounts and newsletters; connect each new story to a course concept

Salary and Career Trajectory: A Data-Grounded View

AI security roles command premium compensation because they require depth in two demanding disciplines simultaneously. The following salary ranges represent 2025 to 2026 market data from Glassdoor, Levels.fyi, and Bureau of Labor Statistics.

Role	Experience Level	US Salary Range (USD)	Growth Trajectory
AI Security Analyst	Entry (0 to 2 years)	\$80,000 to \$105,000	Rapid advancement; high demand for junior talent
AI Risk Manager	Mid (2 to 5 years)	\$110,000 to \$150,000	Strong; regulatory pressure driving demand
AI Governance Engineer	Mid (3 to 6 years)	\$100,000 to \$140,000	Growing rapidly with EU AI Act and US EO implementation
AI Security Architect	Senior (6+ years)	\$130,000 to \$180,000	Very high; few experienced candidates

AI Red Team Lead	Senior (5+ years)	\$140,000 to \$190,000	Extremely high; offensive AI expertise is rare
CISO with AI Specialisation	Executive (10+ years)	\$200,000 to \$350,000+	Peak demand; transformation of security leadership

Course Architecture: How This Learning Experience Is Designed

This course is architected to mirror how working AI security professionals think — not simply to reproduce exam objectives in lecture form. Each section builds on the previous in a deliberate sequence.

- **Foundation first:** The course opens with AI fundamentals because you cannot effectively defend what you do not understand. ML architecture, training dynamics, and model behaviour provide the mental models you need to reason about attack surfaces.
- **Threat-before-defence sequencing:** Within technical sections, attack techniques are taught before defensive countermeasures. Understanding how an adversarial example is constructed before studying defences against it produces deeper retention of both.
- **Framework integration:** Governance frameworks such as NIST AI RMF and EU AI Act are introduced in context throughout technical lectures, not only in the dedicated governance section. Real professionals apply frameworks continuously.
- **Exam objective traceability:** Every lecture maps explicitly to SecAI+ exam objectives. When you complete this course, there are no exam objectives you have not encountered.

Key Terms Glossary

Term	Definition
CompTIA SecAI+	CompTIA's vendor-neutral AI security certification launched in 2026; validates ability to secure AI systems, manage AI risks, and ensure AI compliance
Data Poisoning	Attack that introduces malicious training examples to corrupt an AI model's learned behaviour
Adversarial Example	Input crafted with mathematically optimised perturbations to cause an AI model to misclassify with high confidence
Prompt Injection	Attack where malicious instructions embedded in user input override an LLM's system prompt or intended behaviour
Model Evasion	Category of attacks designed to make malicious inputs appear benign to an AI classifier
Deepfake	AI-generated synthetic media (video, audio, image) that convincingly impersonates a real person
NIST AI RMF	NIST's AI Risk Management Framework; voluntary framework for identifying, measuring, and managing AI risks
EU AI Act	European Union regulation classifying AI systems by risk and imposing requirements on high-risk AI deployments
MLSecOps	Practice of integrating security into the ML development and

	operations lifecycle
AI Red Teaming	Structured adversarial testing of AI systems to identify vulnerabilities before deployment
Domain 3 (38%)	Highest-weighted SecAI+ exam domain covering AI Attack Vectors and Defense
Active Recall	Study technique involving retrieving information from memory rather than re-reading; significantly more effective for retention

Quick Revision Checklist

- I understand why AI security is a distinct discipline from traditional cybersecurity
- I can name at least four types of AI-specific security incidents with real-world examples
- I know the four SecAI+ exam domains, their weightings, and the core competencies each covers
- I understand which domain carries the highest exam weight and why
- I can describe at least three career pathways that SecAI+ opens based on different starting backgrounds
- I know the five active study techniques most effective for technical certification preparation
- I understand the current AI security skills gap and why it creates career opportunity
- I can describe the four learning habits that predict first-attempt certification success
- I have a realistic understanding of salary ranges across AI security roles at different seniority levels
- I understand how this course is structured and why the sequencing of topics is deliberate

20 Exam-Based MCQs — Welcome to the Course

Test yourself after reading. These questions are written in real exam style — scenario-heavy, with plausible distractors. Read every explanation, including for questions you answered correctly.

Q1. Which CompTIA SecAI+ exam domain carries the highest weighting, and what does this weighting reflect about the industry?

- A. Domain 1 (AI Fundamentals) at 38%, reflecting that technical AI knowledge is the hardest to acquire
- B. Domain 3 (AI Attack Vectors and Defense) at 38%, reflecting that understanding offensive AI techniques is central to building effective security
- C. Domain 2 (AI Security Operations) at 38%, reflecting that monitoring is the primary function of AI security teams
- D. Domain 4 (Governance, Risk, and Compliance) at 38%, reflecting the regulatory priority of AI governance

Answer: B **Explanation:** Correct: Domain 3 carries 38% of the exam weight because securing AI systems requires deep understanding of how those systems can be attacked. You cannot effectively defend against adversarial examples, model poisoning, or prompt injection without understanding how these attacks work technically. Domain 1 (14%), Domain 2 (29%), and Domain 4 (19%) are all lower-weighted.

Q2. A machine learning engineer at a bank wants to transition into AI security. Based on typical career pathway timelines, what is the MOST accurate expectation for their exam readiness?

- A. 5 to 8 months, because they need to build both AI and security knowledge from scratch
- B. 2 to 3 months, because they have a strong AI foundation and need to focus on learning security concepts
- C. 6 to 12 months, because AI security requires advanced mathematics beyond what ML engineers typically know
- D. No preparation needed — ML engineers automatically qualify for SecAI+

Answer: B **Explanation:** Correct: ML engineers typically need 2 to 3 months to add security concepts to their existing strong AI foundation. The most efficient path focuses on security operations, attack techniques, and governance frameworks. Career changers from unrelated fields need 5 to 8 months. Advanced mathematics is not a new requirement for ML engineers. The exam must be passed — there is no automatic qualification.

Q3. Which of the following BEST explains why the Samsung ChatGPT incident (2023) was missed by traditional security controls?

- A. Samsung lacked endpoint detection and response software on the affected workstations
- B. The data left the organisation through user-initiated, legitimate API calls — no network intrusion or malware was involved
- C. The incident involved encrypted communication that security tools cannot inspect
- D. Samsung had disabled their data loss prevention tools during the incident period

Answer: B **Explanation:** Correct: In the Samsung incident, employees voluntarily pasted proprietary source code into ChatGPT's interface. From a traditional security perspective, this was an authorised user making authorised API calls — no attack signature, no malware, no intrusion. Traditional DLP tools focused on file exfiltration and email attachments were not designed to detect data being typed into AI chat interfaces. This illustrates why AI-specific security controls are necessary.

Q4. According to ISC2 research cited in the course, what percentage of organisations report lacking sufficient AI security expertise?

- A. 45%
- B. 70%
- C. 78%
- D. 89%

Answer: D **Explanation:** Correct: ISC2 research found that 89% of organisations report lacking sufficient AI security expertise. The 70% figure refers to the percentage of enterprises that have deployed AI in production (McKinsey). The 78% figure refers to organisations that faced AI security incidents in 2024. The 45% figure refers to the annual growth rate of AI security roles on LinkedIn. Each statistic has a specific source and meaning.

Q5. A cybersecurity professional with 4 years of SOC experience wants to study for SecAI+ using the most effective method. Which study approach is MOST evidence-based?

- A. Read the entire official study guide cover-to-cover, then attempt practice tests
- B. Watch all course lectures twice before reviewing any notes or attempting questions
- C. Use daily 45-minute sessions with active recall, spaced repetition, and practice questions from day one
- D. Focus exclusively on Domain 3 since it carries 38% of the exam weight

Answer: C **Explanation:** Correct: Cognitive science research consistently shows that daily spaced practice with active recall significantly outperforms massed study, passive re-reading, or re-watching. Starting practice questions early exploits the testing effect, where retrieval attempts strengthen memory more than additional exposure. Focusing exclusively on Domain 3 would ignore 62% of exam content.

Q6. The Microsoft Tay chatbot incident is MOST relevant to which SecAI+ exam domain?

- A. Domain 4 (Governance, Risk, and Compliance) — because the incident involved a policy failure to filter user inputs
- B. Domain 3 (AI Attack Vectors and Defense) — because users exploited a data poisoning vector by feeding adversarial inputs into the training feedback loop
- C. Domain 2 (AI Security Operations) — because the incident was a monitoring failure that should have been caught by SOC tools
- D. Domain 1 (AI Fundamentals) — because the incident demonstrated weaknesses in transformer architecture

Answer: B **Explanation:** Correct: The Tay incident is a canonical example of feedback loop poisoning, a Domain 3 attack vector. Users discovered that Tay incorporated their inputs into its ongoing learning, then systematically fed it offensive content to poison its outputs. This was an adversarial manipulation of the training mechanism, not a governance policy failure, monitoring gap, or architectural weakness in isolation.

Q7. Which study habit is MOST effective for identifying knowledge gaps before the exam?

- A. Re-reading course notes to reinforce what you already know
- B. Watching lectures a second time at 1.5x speed
- C. Deliberately reviewing incorrect practice question answers in detail and categorising them by domain
- D. Creating a study schedule with time blocked for each domain proportional to exam weighting

Answer: C **Explanation:** Correct: Reviewing incorrect answers in detail is the most targeted gap-identification method. Correct answers confirm existing knowledge; wrong answers map precisely to gaps. Categorising wrong answers by domain enables targeted remediation. Re-reading and re-watching reinforce what you already know rather than addressing gaps. Study scheduling is useful for time management but does not itself identify knowledge gaps.

Q8. An AI security analyst discovers a stop sign near a self-driving car test track has small stickers applied to it, causing the vehicle's vision system to classify it as a speed limit sign. Which AI attack category does this represent?

- A. Data poisoning — the training data was corrupted to misclassify stop signs
- B. Model extraction — the attacker is stealing the vision model's parameters
- C. Adversarial example / model evasion — a physical perturbation causes the deployed model to misclassify at inference time
- D. Prompt injection — text instructions are embedded in the visual input

Answer: C **Explanation:** Correct: Physical adversarial examples exploit the sensitivity of vision models to carefully crafted perturbations. The stickers cause the deployed model to misclassify during inference. They do not affect training data, steal model parameters, or inject text instructions. This is one of the most studied real-world AI security threats because autonomous vehicle safety depends entirely on reliable perception.

Q9. Why does being an early adopter of the SecAI+ certification provide disproportionate career value?

- A. Early adopters receive higher exam scores because the exam is easier when first released
- B. Early certified professionals establish expertise and credibility in a field before credential supply catches up with demand, creating a sustained competitive advantage
- C. CompTIA provides salary bonuses to early adopters of new certifications
- D. The SecAI+ certification expires faster for late adopters, requiring more frequent renewal

Answer: B **Explanation:** Correct: Early adoption creates competitive advantage through two mechanisms: you develop genuine expertise while others are still discovering the field, and when organisations begin requiring the credential, you are already an experienced practitioner. The historical parallel is Security+ professionals who earned it in 2002 to 2004 versus 2015, gaining a decade head start in career development. There are no score differences, salary bonuses, or differential expiry terms.

Q10. The ShadowRay vulnerability (2024) is MOST relevant as an example of which threat category?

- A. Traditional network vulnerability exploitation that happened to affect an AI framework
- B. AI supply chain attack — a vulnerability in an AI-specific infrastructure component (Ray framework) enabled attackers to compromise ML workloads
- C. Social engineering of ML engineers to reveal model parameters
- D. Deepfake-based impersonation of AI system administrators

Answer: B **Explanation:** Correct: ShadowRay represented a supply chain attack specifically targeting AI infrastructure — the Ray distributed computing framework used for training and serving ML models. This is distinct from a traditional network exploit because it targeted AI-specific components and enabled compromise of ML workloads rather than general IT systems. The other options describe different attack categories not relevant to ShadowRay.

Q11. A compliance officer with no technical background is considering SecAI+. Based on the course's stated prerequisites, which statement is MOST accurate?

- A. They cannot sit the exam because they lack the required Security+ prerequisite
- B. They are ineligible because SecAI+ requires 5 years of IT experience as a mandatory prerequisite
- C. They can succeed with additional study — SecAI+ has no mandatory prerequisites, and the course covers necessary fundamentals
- D. They should wait until they have completed a machine learning degree before attempting the certification

Answer: C **Explanation:** Correct: CompTIA recommends experience but does not mandate prerequisites for SecAI+. Motivated learners from non-technical backgrounds can succeed with structured study. The course is designed to cover AI fundamentals for those without prior ML knowledge. Security+ is recommended but not required. The 3 to 5 years experience is a recommendation, not a mandatory requirement.

Q12. Which of the following BEST describes why the teach-back method is effective for SecAI+ preparation?

- A. Teaching others earns Continuing Education Units that count toward certification renewal
- B. Attempting to explain a concept in plain language reveals specific gaps in your own understanding before the exam does
- C. Teaching peers allows you to share study materials legally
- D. The method reduces the number of practice questions required to pass

Answer: B **Explanation:** Correct: The teach-back method works because fluent, accurate explanation of a concept requires genuine understanding. Gaps in your explanation reveal gaps in your knowledge in a low-stakes setting. This is cognitively distinct from recognising correct answers on practice questions, where familiarity can masquerade as understanding.

Q13. A student plans to study for SecAI+ with an 8-hour intensive session every Sunday rather than daily 30-minute sessions. What is the MOST likely outcome compared to daily practice?

- A. The Sunday-only approach will be more effective because longer uninterrupted sessions allow deeper focus
- B. Both approaches are equivalent as long as the total weekly hours are the same
- C. The Sunday-only approach will produce lower retention because sleep-based memory consolidation and spaced repetition are eliminated from six days of the week
- D. The Sunday-only approach is preferable for technical content because complex topics require sustained cognitive immersion

Answer: C **Explanation:** Correct: Memory consolidation occurs primarily during sleep. Daily study allows the brain to consolidate the previous day's learning each night, creating stronger and more retrievable memory traces. Massed practice produces rapid short-term retention but significantly faster forgetting. Technical certification content is particularly susceptible to this effect because concepts are interrelated and require multiple retrieval events to fully encode.

Q14. Which salary range BEST represents a senior AI Security Architect with 7 years of experience in a major US technology hub in 2025 to 2026?

- A. \$80,000 to \$105,000
- B. \$100,000 to \$140,000
- C. \$130,000 to \$180,000 base, potentially higher with geographic premium
- D. \$50,000 to \$75,000

Answer: C **Explanation:** Correct: AI Security Architects are senior-level roles typically requiring 6 or more years of experience, with base salary ranges of \$130,000 to \$180,000 in the US market. In major technology hubs, this figure can increase by 25 to 40%, potentially reaching \$160,000 to \$250,000+ total compensation. The lower ranges correspond to analyst and mid-level roles. Option D falls well below market rates for this seniority and specialisation.

Q15. The course sequences attack techniques before defensive countermeasures within each technical section. What is the pedagogical rationale for this design choice?

- A. Exams always test offensive knowledge before defensive knowledge, so the sequencing mirrors exam structure
- B. Understanding how an attack is constructed provides the mental model needed to understand why specific countermeasures work, producing deeper retention of both
- C. Offensive techniques are simpler than defensive techniques and should always be taught first
- D. CompTIA requires all certified training materials to cover attacks before defences

Answer: B **Explanation:** Correct: Learning why a defence works requires understanding the attack it defends against. Teaching defences without attack context produces surface-level memorisation — students know what to do but not why. Teaching the attack first creates the cognitive need that the defence satisfies, a well-validated instructional design principle.

Q16. Domain 4 (Governance, Risk, and Compliance) accounts for 19% of the SecAI+ exam. Which regulatory development MOST explains growing industry demand for this competency?

- A. The introduction of GDPR in 2018, which first required data protection officers
- B. The EU AI Act and US Executive Order on AI, which impose specific AI risk assessment and documentation requirements on organisations deploying AI systems
- C. The expansion of PCI-DSS to cover AI-based payment fraud detection systems
- D. The deprecation of the ISO 27001 standard in favour of AI-specific frameworks

Answer: B **Explanation:** Correct: The EU AI Act (2024) and US Executive Order on Safe, Secure, and Trustworthy AI (2023) are the primary regulatory drivers creating demand for AI governance expertise. These regulations impose specific risk classification, documentation, human oversight, and conformity assessment requirements on organisations deploying AI. GDPR predates AI-specific regulation. PCI-DSS has not been specifically expanded to cover AI fraud detection. ISO 27001 has not been deprecated.

Q17. A junior analyst argues that since they already have Security+, they can skip Domain 1 (AI Fundamentals) of their SecAI+ study. Why is this reasoning flawed?

- A. Security+ and SecAI+ use different multiple-choice question formats
- B. Domain 1 covers AI and ML concepts such as model training, neural network architectures, and statistical learning that are not covered in Security+ and form the prerequisite knowledge for Domains 2 and 3
- C. CompTIA rules prohibit candidates with Security+ from taking shortcuts in SecAI+ preparation
- D. Domain 1 accounts for 38% of the exam, making it too risky to skip

Answer: B **Explanation:** Correct: Security+ covers traditional security operations, network security, and threat intelligence — it does not cover ML algorithms, neural network architectures, training dynamics, or statistical learning fundamentals. These AI-specific concepts are tested in Domain 1 and are the prerequisite knowledge required to understand AI-specific attack techniques in Domain 3. Option D has the domain weighting wrong — Domain 1 is 14%, not 38%.

Q18. Which of the following BEST illustrates the concept mapping study technique applied to SecAI+ preparation?

- A. Creating a table that lists each exam domain with its percentage weighting
- B. Drawing a diagram showing how adversarial examples relate to ML model architecture, which relates to training data quality, which relates to data governance policies
- C. Writing out exam objective lists in alphabetical order for easier memorisation
- D. Highlighting key terms in the official study guide in different colours for each domain

Answer: B **Explanation:** Correct: Concept mapping involves creating visual representations of the relationships between concepts, not just lists or tables. Option B captures the interconnected nature of AI security knowledge — an adversarial example exploits properties of the model architecture that emerged from the training process, which depends on data quality, which is controlled by governance policies. Options A, C, and D are organisational techniques, not concept mapping.

Q19. What is the primary reason AI security incidents like the Arup Engineering deepfake fraud are difficult to prevent using traditional cybersecurity controls?

- A. The attackers used zero-day exploits that no antivirus could detect
- B. The attack exploited human trust in visual and audio confirmation — a non-technical vulnerability that traditional security tools are not designed to address
- C. The organisation's network lacked sufficient encryption on video conferencing channels
- D. The attack originated from a nation-state actor with resources beyond most organisations' defences

Answer: B **Explanation:** Correct: The Arup incident involved no malware, no network intrusion, and no technical vulnerability exploitation. The attack was entirely social engineering executed through AI-generated media that exploited the human tendency to trust what they see and hear. Traditional security controls do not protect against this vector. Effective defence requires procedural controls such as out-of-band verification and multi-party approval workflows.

Q20. A student scored 65% on their first full practice exam with 3 weeks until their exam date. Which study plan is MOST likely to improve their score to a passing level?

- A. Re-watch all course lectures from the beginning to refresh the content
- B. Categorise every wrong practice answer by domain, focus intensive study on the weakest domains using active recall, retake a different practice exam after one week, and repeat the cycle
- C. Read the official CompTIA study guide cover-to-cover once
- D. Increase daily study sessions to 4 hours of continuous reading

Answer: B **Explanation:** Correct: With 3 weeks and a 65% baseline, the highest-leverage approach is targeted gap remediation. Categorising wrong answers identifies exactly which domains and objectives need work. Active recall study of those specific topics is more efficient than reviewing everything. Re-testing after a week measures progress and identifies remaining gaps. Re-watching lectures or reading the study guide without targeting specific gaps is inefficient. Extended passive reading sessions do not leverage spaced repetition or active recall.