

Understanding Artificial Intelligence Fundamentals

CompTIA SecAI+ | Section 2: Introduction to AI Types and Techniques

1. Rethinking What a Computer Does

Every computer program ever written before AI followed the same basic contract: you tell it exactly what to do, and it does exactly that — nothing more, nothing less. Click Save, the file saves. Enter the wrong password, access is denied. This model — called deterministic computing — is extraordinarily reliable, but it has a ceiling. It can only handle situations the programmer anticipated. Anything outside the rulebook produces failure or silence.

Artificial intelligence breaks this contract on purpose. An AI system is not given a complete rulebook. Instead, it is given data — thousands or millions of examples — and it figures out the rules itself. The result is a system that can handle situations nobody programmed it for, because it learned a model of the world rather than a set of explicit instructions.

This distinction is not a technical detail. It is the reason AI is genuinely transformative in security rather than just another tool.

Key Point: Traditional computing is deterministic — the same input always produces the same output based on programmed rules. AI is probabilistic — it produces predictions with confidence levels based on learned patterns. This shift from rules to patterns is the foundation of every AI security application.

2. How AI Actually Learns: Patterns From Examples

The spam email analogy from the lecture is worth unpacking more carefully, because it captures the entire mechanism of machine learning in miniature.

Nobody gave you a written list of every spam phrase. You were never handed a rulebook that said "if the email contains 'Nigerian prince' AND 'wire transfer' AND exclamation marks, classify as spam." Instead, you saw hundreds of spam emails over years, your brain unconsciously extracted the statistical patterns, and now you recognise spam instantly — even when the spammer invents a new phrase you have never seen before. You generalised from examples.

AI does precisely this, but at industrial scale. A spam filter trained on ten million emails has seen far more examples than any human, and it has extracted patterns far more subtle than anything a human analyst could articulate — slight variations in sending domains, header anomalies, structural patterns in URLs, combinations of signals that individually look harmless but together predict malice.

Why Generalisation Matters More Than Memorisation

A critical distinction: the goal of AI training is generalisation, not memorisation. A model that simply memorised its training data would fail the moment it encountered anything new. The model must learn patterns abstract enough to apply to unseen data. This is why AI can detect a brand-new malware variant it has never seen — not because it recognises the file, but because it recognises the behavioural pattern of malware in general.

3. The Signature-Based Detection Model and Its Limits

To appreciate why AI matters in security, you need to understand what came before it — and why it worked so well for so long before breaking down.

Signature-based detection treats every known threat as a unique fingerprint. A virus, a piece of malware, or a malicious network packet has certain byte sequences, checksums, or behavioural markers that distinguish it from legitimate software. Security vendors catalogue these fingerprints into databases. When your antivirus scans a file, it compares the file against every known fingerprint. If it finds a match, it raises an alert.

Dimension	Signature-Based Detection	AI-Based Detection
What it recognises	Exact matches to known threats	Patterns associated with malicious behaviour
Response to new threats	Blind — no signature exists yet	Can flag anomalies even without prior examples
Update dependency	Requires constant signature database updates	Improves automatically with new training data
False positive risk	Low for known threats	Higher — probabilistic, requires tuning
Computational cost	Very low — simple lookup	Higher — model inference at every event
Evasion difficulty	Easy — change a few bytes	Hard — must change the entire behavioural pattern

The Zero-Day Problem

A zero-day vulnerability is one the vendor does not yet know about. A zero-day exploit weaponises it. Because no signature exists for an attack that has never been seen, signature-based systems offer exactly zero protection. The attacker has the initiative: they craft the attack, launch it, and the defender's tools remain blind until a researcher discovers the attack, reverse-engineers it, creates a signature, and ships an update — a cycle that can take days to weeks. During that window, the attack runs unchecked.

Attackers industrialised this gap. Modern malware factories produce thousands of signature variants per day — minor mutations that preserve malicious function while evading exact-match detection. The defender's signature database is always chasing an adversary who moves faster.

4. How AI Inverts the Detection Problem

Signature-based detection starts with the question: "Do I recognise this as bad?" This is a blacklist approach. AI-based detection starts with the opposite question: "Does this fit the pattern of normal behaviour?" This is a baseline approach — and it is far more powerful against novel threats.

An AI security model spends its training phase learning what normal looks like. Normal login times for each user. Normal data transfer volumes. Normal sequences of system calls for legitimate software. Normal network destinations for each application. Once the baseline is established, any meaningful

deviation triggers investigation — not because the system recognises the threat, but because it recognises the anomaly.

This means a genuinely novel attack — one nobody has ever seen — can still be detected, because novel malicious behaviour almost always deviates from normal behaviour in ways the model was trained to notice. The attacker must now evade not a list of known bad things, but a statistical model of everything that looks normal. That is a much harder problem.

Real-World Incident — SolarWinds Supply Chain Attack (2020): The SolarWinds Orion compromise was one of the most sophisticated supply chain attacks in history. Attackers embedded malicious code (SUNBURST) into legitimate software updates distributed to approximately 18,000 organisations including US government agencies. Traditional signature-based tools missed it entirely because the malicious code was signed with SolarWinds' own digital certificate — it looked legitimate. AI-based anomaly detection systems at some affected organisations flagged unusual outbound traffic patterns to suspicious domains weeks before investigators identified the campaign. The incident demonstrated that when an attacker achieves trusted access, only behavioural anomaly detection — the kind AI enables — stands between them and complete invisibility.

5. The Alert Volume Crisis

Even if traditional security tools could detect every threat, they would face a second problem: the sheer volume of data modern infrastructure generates. A mid-sized enterprise with five thousand employees generates millions of security-relevant events every day — logins, file accesses, network connections, process executions, configuration changes. A large enterprise generates billions.

Security operations centres (SOCs) at large organisations routinely receive ten thousand or more alerts per day. Studies consistently show that analysts review only a fraction of these alerts, and that alert fatigue — the psychological state where analysts stop taking alerts seriously because most are false positives — is a primary cause of missed detections in real breaches.

AI as a Force Multiplier, Not a Replacement

AI addresses this crisis not by replacing human analysts but by doing the work that does not require human judgment: ingesting raw event data at machine speed, correlating signals across disparate systems, filtering out noise, and surfacing only the small percentage of alerts that genuinely warrant human review. A well-configured AI triage layer can reduce analyst workload by 80–90% while simultaneously improving detection rates — because it catches patterns across thousands of systems that no individual analyst could hold in their head.

The human analyst then works with higher-quality information: pre-correlated, prioritised, and often accompanied by an AI-generated summary of why this activity was flagged. This is the practical reality of how AI integrates into security operations today.

6. Deterministic vs Probabilistic Systems

Understanding the difference between deterministic and probabilistic systems is essential for the SecAI+ exam and for practical work in AI security.

Property	Deterministic (Traditional)	Probabilistic (AI)
Decision basis	Explicit rules written by humans	Statistical patterns learned from data
Output type	Binary: yes/no, allow/deny	Confidence score: 0–100% probability
Handling novel input	Fails or defaults to no-match	Produces best estimate with confidence level
Explainability	Fully explainable — rule chain is traceable	Often partially opaque (depends on model type)
Failure mode	Blind to anything outside the rules	Can produce confident wrong answers (hallucination)
Example in security	Firewall rule: block port 22 from external IPs	UEBA: 87% probability this login is credential stuffing

The probabilistic nature of AI outputs has a practical implication for security workflows: AI systems should almost never make final binary decisions alone. They should produce scored recommendations that human analysts or downstream rule-based systems act upon. This is why the SecAI+ exam emphasises human oversight, explainability, and responsible AI — probabilistic systems can be wrong in ways deterministic systems cannot.

7. Three Core AI Capabilities in Cybersecurity

The lecture identifies three capabilities AI brings to security. Understanding what each capability means at a technical level — not just as a label — is what exam questions will test.

Capability	What It Means	Security Example
Pattern Recognition	Identifying known or anomalous signatures in data by comparing against a learned baseline or classification model	Network intrusion detection: flagging packets whose protocol behaviour matches learned attack patterns
Prediction	Forecasting future states or events by extrapolating from historical patterns — assigning probabilities to outcomes that have not yet occurred	Vulnerability prioritisation: predicting which unpatched CVEs are most likely to be exploited in the next 30 days based on exploit-kit activity and asset exposure
Automation	Taking autonomous or semi-autonomous action based on AI-generated decisions, without waiting for human input	SOAR integration: automatically isolating an endpoint when an AI model assigns >90% malware probability to a detected process

These three capabilities are not independent. A mature AI security system uses all three in sequence: pattern recognition identifies the anomaly, prediction scores its severity and likely next steps, and automation executes a proportionate response — all within seconds of the initial event.

8. The Dual-Use Reality: AI Is a Tool for Both Sides

One of the most important conceptual shifts in the SecAI+ certification is the explicit acknowledgement that AI is a dual-use technology. The same capabilities that make AI powerful for defenders make it equally powerful for attackers.

How Defenders Use AI

- Anomaly detection: identifying deviations from baseline behaviour at machine speed across millions of events
- Malware analysis: automated classification of unknown files using behavioural models rather than signatures
- User and Entity Behaviour Analytics (UEBA): detecting insider threats and compromised credentials by modelling individual user behaviour
- Threat intelligence: correlating indicators across open-source and commercial feeds automatically
- SOAR automation: executing playbook responses faster and more consistently than human analysts

How Attackers Use AI

- Spear-phishing at scale: generating personalised, grammatically correct phishing emails from scraped social media data — removing the spelling errors that trained users to recognise phishing
- Adversarial attacks: crafting inputs specifically designed to fool AI security models — making malware that looks normal to an anomaly detector
- AI-powered reconnaissance: using AI to rapidly enumerate attack surfaces, identify unpatched systems, and prioritise targets
- Deepfakes: synthesising voice or video of executives to authorise fraudulent transfers or bypass voice authentication
- Automated vulnerability discovery: using AI to find novel code flaws faster than human researchers

Key Point: The SecAI+ exam expects you to reason about AI from both perspectives. A question that describes an AI-powered capability may be asking about a defensive tool or an offensive weapon — context determines the answer. Never assume a capability is inherently defensive or offensive.

9. Key Terms Glossary

Term	Definition
Artificial Intelligence (AI)	A system that learns from data to perform tasks that traditionally required human intelligence, including pattern recognition, prediction, and decision-making
Deterministic system	A computing system where the same input always produces the same output, based on explicitly programmed rules
Probabilistic system	A system that produces outputs as confidence scores or probability distributions rather than binary decisions
Machine learning (ML)	A subset of AI where systems improve their performance on tasks through exposure to training data, without being explicitly reprogrammed
Signature-based detection	Security technique that identifies threats by matching observed data against a database of known threat fingerprints

Zero-day	A vulnerability or exploit that is unknown to the vendor, for which no patch or signature exists
Anomaly detection	AI technique that establishes a baseline of normal behaviour and flags statistically significant deviations
Generalisation	The ability of a trained AI model to perform correctly on data it has not previously seen
Alert fatigue	The state in which security analysts become desensitised to alerts due to excessive volume, leading to missed detections
Dual-use technology	Technology that has both legitimate defensive and potentially offensive applications
UEBA	User and Entity Behaviour Analytics — AI-driven analysis of human and system behaviour to detect insider threats and compromised accounts
SOAR	Security Orchestration, Automation and Response — platforms that use AI to automate security workflows and incident response

10. Quick Revision Checklist

- AI learns from data examples rather than following explicit programmer-defined rules
- Traditional computing is deterministic; AI is probabilistic — outputs are confidence scores, not binary certainties
- Signature-based detection matches against known threat fingerprints and is blind to zero-days
- Attackers evade signatures by mutating malware; AI-based anomaly detection defeats this by modelling normal behaviour
- The zero-day problem: no patch and no signature exist; behavioural AI is the primary detection mechanism
- Alert fatigue is a real operational crisis — AI addresses it by triaging and prioritising at machine speed
- Three AI capabilities in security: pattern recognition, prediction, automation — these work sequentially in mature systems
- AI is deterministic in its architecture but probabilistic in its outputs — it produces scores, not rules
- The dual-use nature of AI means the same capability can detect phishing or generate it — exam questions exploit this ambiguity
- Generalisation (performing on unseen data) is the goal of training — not memorisation
- AI does not replace human analysts; it acts as a force multiplier by filtering and prioritising

20 Exam-Based MCQs — AI Fundamentals for Security

Test yourself after reading. These questions are written in real exam style — scenario-heavy, with plausible distractors. Read every explanation, including for questions you answered correctly.

Q1. What is the fundamental difference between a traditional rule-based security system and an AI-based security system?

- A) Traditional systems are faster; AI systems are more accurate

- B) Traditional systems follow explicitly programmed rules; AI systems learn patterns from data
- C) Traditional systems use cloud infrastructure; AI systems run on-premises
- D) Traditional systems require more storage; AI systems are more computationally efficient

Answer: B **Explanation:** B is correct because the defining characteristic of AI is that it derives its decision-making logic from patterns in training data rather than explicit human-written rules. A is wrong — AI systems are not inherently more accurate than all rule-based systems, and speed depends on implementation. C is wrong — both deployment models exist for both types of system. D is wrong — AI model inference is typically more computationally expensive than a signature lookup.

Q2. A security vendor's antivirus product failed to detect a brand-new ransomware variant, even though the ransomware performed the same file-encryption behaviour as dozens of known variants. What limitation does this BEST illustrate?

- A) The antivirus product had an outdated licence
- B) The ransomware used AI-generated code
- C) Signature-based detection cannot detect threats it has not previously catalogued
- D) The antivirus product lacked sufficient computing resources

Answer: C **Explanation:** C is correct — signature-based detection requires an exact or near-exact match to a known fingerprint. A new ransomware variant, even one with identical behaviour, has a different binary fingerprint and evades detection. A is wrong — licence status does not affect detection capability. B is wrong — the question does not state AI was used, and even if it were, the root issue is the signature model's limitation. D is wrong — resource constraints do not cause missed detections in signature matching.

Q3. Scenario: A financial services company deploys an AI-based User and Entity Behaviour Analytics (UEBA) platform. On a Tuesday morning, the system flags a senior finance manager's account with a 94% probability anomaly score. The manager is logging in from her usual office location at her usual time, but the system has flagged that she accessed 4,200 files in 45 minutes — a volume that would normally take her two days. What does this scenario BEST demonstrate about AI-based security detection?

- A) The AI system is producing a false positive because the login location and time are normal
- B) AI systems require human-written rules to flag specific file access thresholds
- C) AI detected a behavioural anomaly that no signature or rule would have caught
- D) The UEBA platform should be reconfigured to ignore file access volume as a signal

Answer: C **Explanation:** C is correct — the AI model established a baseline of the manager's normal file access behaviour and flagged a statistically significant deviation (4,200 files in 45 minutes vs her normal pattern) even though the login itself appeared legitimate. This is exactly the kind of anomaly that signature-based tools — which look for known-bad patterns — would miss entirely. A is wrong — a normal login time and location do not invalidate an anomalous behavioural pattern; credential stuffing attacks frequently originate from expected IP addresses after initial compromise. B is wrong — the AI learned the threshold from data, not from a human-written rule. D is wrong — file access volume is a high-value signal for detecting data exfiltration and insider threats.

Q4. Which of the following BEST describes a zero-day attack?

- A) An attack that begins exactly at midnight
- B) An attack that exploits a vulnerability unknown to the vendor, for which no signature or patch exists
- C) An attack that targets newly installed systems with default configurations
- D) An attack that overwrites the master boot record on day zero of deployment

Answer: B **Explanation:** B is the correct definition — a zero-day refers to the vendor having zero days of knowledge of the vulnerability, meaning no defensive signature or patch has been published. A is wrong — the name refers to vendor response time, not the time of day the attack occurs. C is wrong — attacks on default-configured systems are a different category (misconfigurations) and are not zero-days unless they exploit an unknown vulnerability. D is wrong — boot record attacks are a specific technique, not a definition of zero-day.

Q5. *Scenario:* A SOC analyst at a large retail company receives 12,000 security alerts on a typical Monday. After filtering, the team manually investigates approximately 200. A recent audit found that three confirmed breaches over the past year were missed because the relevant alerts were buried in the queue. The CISO is evaluating AI-based triage tools. Which AI capability would MOST directly address the SOC's core operational problem?

- A) Automation — automatically blocking all flagged IP addresses without human review
- B) Prediction — forecasting which assets are most likely to be attacked in the next quarter
- C) Pattern recognition — identifying and prioritising the small percentage of alerts that warrant human investigation
- D) Replacing human analysts with AI systems entirely to eliminate the volume problem

Answer: C **Explanation:** C is correct — the core problem is alert fatigue from volume. AI pattern recognition that triages, scores, and prioritises alerts allows the team to focus on the ~1–5% of alerts that represent genuine threats, which is how the three missed breaches would have been surfaced. A is wrong — automatically blocking all flagged IPs would produce an unacceptable false-positive rate and disrupt legitimate business traffic. B is wrong — prediction helps with vulnerability management but does not address the immediate alert triage problem. D is wrong — AI augments analysts; fully autonomous decision-making at the blocking level introduces unacceptable risk of business disruption.

Q6. An AI security model produces an output of '82% probability — suspicious login.' A traditional firewall rule produces an output of 'BLOCKED.' Which statement BEST characterises the difference in output type?

- A) The AI model is less reliable because it is not certain
- B) The firewall is more sophisticated because it makes a definitive decision
- C) The AI model produces probabilistic output; the firewall produces deterministic output
- D) Both systems are deterministic but the AI model uses a different decision threshold

Answer: C **Explanation:** C correctly identifies the core distinction — AI produces probabilistic scores representing confidence levels, while deterministic systems produce definitive binary outputs based on programmed rules. A is wrong — probabilistic output is not a reliability defect; it provides richer information that human analysts can act on. The 18% uncertainty can inform investigation depth. B is wrong — definitiveness is not sophistication; a firewall blocking a legitimate business partner's IP is a definitive failure. D is wrong — the AI model is fundamentally probabilistic in architecture, not a deterministic system with a different threshold.

Q7. Scenario: *A threat actor group has developed a spear-phishing campaign targeting CFOs at manufacturing companies. Unlike previous campaigns, the phishing emails are personalised with accurate details about each target's recent business activities, are grammatically perfect, reference real suppliers by name, and arrive at times when the target is known to be checking email. Researchers suspect AI was used to generate the emails. What does this scenario MOST likely illustrate about the role of AI in offensive operations?*

- A) AI enables attackers to automate and personalise social engineering at scale, removing previously reliable human error indicators
- B) AI is only useful to defenders; offensive AI use requires nation-state resources
- C) The attack demonstrates a signature-based detection failure, not AI involvement
- D) Personalised phishing is easily detected by AI-based email security tools

Answer: A **Explanation:** A is correct — one of the most impactful offensive uses of AI is generating highly convincing spear-phishing content at scale. Historically, phishing was detectable by poor grammar, implausible details, and generic content. AI-generated phishing eliminates these tells, dramatically increasing success rates. B is wrong — AI tools capable of generating convincing phishing content are widely accessible; this is not limited to nation-state actors. C is wrong — the question explicitly attributes the sophistication to likely AI use; the scenario is illustrating dual-use AI risk, not a detection system failure. D is wrong — AI email security tools are improving but personalised, contextually accurate phishing remains a significant detection challenge.

Q8. In the context of AI security systems, what does 'generalisation' mean and why is it important?

- A) The ability of a model to memorise all training examples and reproduce them exactly
- B) The ability of a model to perform correctly on data it was not trained on, enabling detection of novel threats
- C) The process of converting a specialised AI model into a general-purpose language model
- D) The technique of applying security rules broadly across all systems rather than targeted deployments

Answer: B **Explanation:** B is correct — generalisation is the property that makes AI useful for security. A model that only recognised threats it had seen in training would be equivalent to an advanced signature database. Generalisation allows the model to apply learned patterns to entirely new threats. A is wrong — this describes overfitting or memorisation, which is actually a failure mode. An overfit model performs poorly on new data. C is wrong — this describes a different concept (model conversion/transfer learning) unrelated to generalisation as a learning property. D is wrong — this describes a deployment strategy, not

an AI learning property.

Q9. Which statement BEST describes how AI-based anomaly detection avoids the zero-day problem that defeats signature-based systems?

- A) AI systems maintain larger and more frequently updated signature databases
- B) AI systems require internet connectivity to check threats against real-time cloud databases
- C) AI establishes a baseline of normal behaviour and flags deviations, regardless of whether the specific threat has been seen before
- D) AI systems use cryptographic hashing to detect modified malware variants

Answer: C **Explanation:** C is correct — anomaly detection sidesteps the zero-day problem by inverting the question. Rather than asking 'do I recognise this as bad?', it asks 'does this deviate from normal?' A completely novel attack that has never been documented anywhere will still cause anomalies in system behaviour, network traffic, or user activity. A is wrong — larger signature databases still cannot contain signatures for unknown threats. B is wrong — cloud lookups still rely on known-bad databases; they do not help with truly novel threats. D is wrong — hash-based detection is a form of signature detection and has the same zero-day blind spot.

Q10. A malware author generates 5,000 unique variants of the same ransomware by automatically modifying non-functional byte sequences. All variants perform identical encryption behaviour. What detection approach would MOST effectively catch all 5,000 variants?

- A) Signature-based detection with a comprehensive hash database updated daily
- B) Behavioural AI that detects the encryption activity pattern regardless of file hash
- C) Network perimeter firewall rules blocking known ransomware command-and-control domains
- D) User security awareness training to prevent employees from executing unknown files

Answer: B **Explanation:** B is correct — behavioural AI identifies the ransomware activity pattern (rapid file enumeration, encryption of documents, deletion of shadow copies) which is consistent across all 5,000 variants regardless of their different byte sequences. A is wrong — 5,000 unique hashes require 5,000 unique signatures. If each variant is new, none will be in the database until discovered and catalogued, by which time the attack may be complete. C is wrong — C2 domain blocking helps contain an active infection but does not prevent initial execution; also, novel C2 infrastructure would not be in block lists. D is wrong — user training reduces the probability of execution but cannot detect or stop all executions; it is a complementary control, not a detection mechanism.

Q11. Which of the following is an example of AI being used for prediction in a cybersecurity context?

- A) An AI model that classifies incoming network packets as malicious or benign in real time
- B) An AI model that forecasts which unpatched vulnerabilities in an organisation's estate are most likely to be exploited in the next 30 days
- C) An AI model that automatically quarantines endpoints identified as compromised
- D) An AI model that generates summaries of security incident reports

Answer: B **Explanation:** B is correct — prediction involves forecasting future states or probabilities. Prioritising unpatched CVEs by exploitation likelihood is a prediction task: the model estimates future attacker behaviour based on current threat intelligence and historical exploitation patterns. A is wrong — real-time classification is pattern recognition. C is wrong — automatic quarantine is automation. D is wrong — summarisation is a natural language processing task that could be called pattern recognition or information extraction, but it is not prediction of a future security event.

Q12. Alert fatigue in a security operations centre is BEST described as:

- A) Physical exhaustion caused by 24/7 shift work
- B) The tendency of analysts to become desensitised to alerts due to excessive volume, increasing the risk of missed detections
- C) The degradation of AI model performance when alert volume exceeds processing capacity
- D) Regulatory fatigue caused by too many compliance alert requirements

Answer: B **Explanation:** B is the correct definition — alert fatigue is a psychological and operational phenomenon where analysts receive so many alerts that they begin dismissing them without adequate investigation, treating alerts as background noise. This has been identified as a contributing factor in several high-profile breaches. A is wrong — physical fatigue from shift work is a separate issue. C is wrong — alert fatigue is a human phenomenon, not a system performance issue. D is wrong — regulatory compliance is a separate domain; this term has a specific meaning in security operations.

Q13. *Scenario:* During a post-incident review, a company discovers that their AI-based SIEM flagged the initial compromise with a 78% anomaly score six hours before the breach was detected by a human analyst. The alert was not investigated because analysts were working through a backlog of 8,000 other alerts, most of which were false positives. The company is now evaluating what additional controls to add. Which recommendation would MOST directly address the root cause of the delayed detection?

- A) Replace the AI system with a signature-based tool that produces only confirmed detections
- B) Implement AI-based alert prioritisation and triage to surface high-confidence anomalies for immediate review
- C) Hire additional analysts to process the full alert backlog manually
- D) Increase the anomaly detection threshold to 95% to reduce the number of alerts generated

Answer: B **Explanation:** B is correct — the root cause is that a high-quality alert (78% confidence initial compromise) was buried in an undifferentiated queue of 8,000 alerts. AI-based prioritisation would rank that alert at the top, ensuring it received immediate attention. A is wrong — signature-based tools would have produced zero alerts for this initial compromise (which was presumably a novel or low-indicator event that triggered the anomaly model); they would make the situation worse. C is wrong — adding analysts without changing the alert volume creates a cost-intensive partial solution; the queue will regenerate. D is wrong — raising the threshold to 95% would likely exclude the initial compromise alert (78%) entirely, guaranteeing the breach is not detected early.

Q14. Which of the following statements about AI and cybersecurity attackers is TRUE?

- A) Attackers cannot use AI because AI tools require vendor agreements that exclude malicious use
- B) AI tools are exclusively used by well-funded nation-state threat actors
- C) Attackers use AI to enhance phishing, automate reconnaissance, generate malware, and evade AI-based defences
- D) AI makes attackers detectable because AI-generated content always contains statistical artefacts

Answer: C **Explanation:** C is correct — AI is a dual-use technology. Attackers leverage AI for personalised spear-phishing, automated target discovery, adversarial examples that fool detection models, and rapid vulnerability identification. A is wrong — AI tools including large language models are widely accessible through open-source projects, dark web marketplaces, and standard consumer products; vendor agreements do not prevent malicious use. B is wrong — commodity AI tools have democratised sophisticated capabilities to criminal groups at all levels. D is wrong — while AI-generated content detection research is active, no reliable universal detection method exists, and sophisticated actors can post-process output to reduce detectable artefacts.

Q15. The SOAR acronym stands for Security Orchestration, Automation and Response. In an AI-integrated security context, which function does AI PRIMARILY enhance?

- A) Orchestration — connecting and coordinating different security tools
- B) Automation — executing playbook actions based on AI-generated threat scores without waiting for human authorisation
- C) Response — filing incident reports and communicating with affected stakeholders
- D) Reporting — generating regulatory compliance documentation

Answer: B **Explanation:** B is correct — AI enhances the automation function of SOAR most directly. AI generates threat scores and classifications that trigger automated playbook execution: isolating endpoints, blocking IPs, revoking tokens, or opening incident tickets — without requiring a human to review and authorise each action. This dramatically reduces mean time to respond (MTTR). A is wrong — orchestration (connecting tools) is primarily an integration engineering function; AI contributes to deciding what to do with integrated data, not the plumbing itself. C is wrong — stakeholder communication is a human-in-the-loop process. D is wrong — compliance reporting is a structured data and documentation function.

Q16. A security AI model was trained on six months of network traffic data. Six months after deployment, its detection rate for certain attack types has declined. What AI concept MOST likely explains this?

- A) The model has been updated with new signatures that conflict with the original training
- B) Model drift — the real-world data distribution has changed since training, reducing model accuracy
- C) The model has reached its computational capacity and is processing events too slowly
- D) Alert fatigue has reduced the model's confidence threshold settings

Answer: B **Explanation:** B is correct — model drift (also called data drift or concept drift) occurs when the statistical distribution of real-world data shifts away from the training data distribution. Networks evolve: new applications, protocols, user behaviours, and attack techniques appear that were not

represented in the training data. The model's learned patterns become progressively less accurate. A is wrong — ML models are not signature-based and do not have signature conflicts. C is wrong — computational capacity is an infrastructure issue that would produce latency, not reduced detection accuracy. D is wrong — alert fatigue is a human operational problem; it does not affect model confidence settings.

Q17. Scenario: A cybersecurity team is evaluating two products: Product A is a next-generation antivirus that matches files against a continuously updated cloud hash database refreshed every four hours. Product B is an endpoint AI that monitors process behaviour in real time, comparing activity against a baseline built from 90 days of normal operation. A new malware family is released by a threat actor group. No signatures exist yet. In the first four hours after the malware is released, which product provides BETTER protection against this specific threat, and why?

- A) Product A, because cloud databases receive threat intelligence updates faster than local AI models
- B) Product A, because hash-based detection is more accurate than behavioural detection
- C) Product B, because behavioural AI can flag the malware's anomalous process activity even without a signature
- D) Both products provide equal protection because all modern security tools share threat intelligence in real time

Answer: C Explanation: C is correct — in the zero-day window before any signature exists, Product A is completely blind. The cloud database is updated every four hours, and during those four hours there is no hash to match. Product B detects anomalous process behaviour (unusual memory access patterns, network connections, file operations) that deviate from the established 90-day baseline — regardless of the malware's hash. A is wrong — even a four-hour update cycle creates a four-hour blind window, and new malware may not be included until after discovery, analysis, and cataloguing, which can take longer. B is wrong — hash accuracy is irrelevant when the hash does not exist in the database. D is wrong — real-time threat intelligence sharing does not produce a signature instantly; discovery and cataloguing take time.

Q18. Which of the following BEST describes why AI is described as a 'force multiplier' for security analysts rather than a replacement?

- A) AI is too expensive to fully replace human analysts in most organisations
- B) AI handles the high-volume, pattern-based triage work, freeing analysts to apply judgment, context, and expertise to investigated cases
- C) Regulatory requirements in most jurisdictions mandate that a human make all security decisions
- D) AI systems require human operators to function and cannot run unsupervised

Answer: B Explanation: B is correct — the force multiplier concept means that AI allows a fixed number of analysts to investigate more threats by handling the work that does not require human judgment: ingesting raw data, correlating signals, filtering noise, and scoring alerts. Human analysts then apply contextual expertise to the cases AI surfaces. The total output of the human+AI team is greater than either could achieve alone. A is wrong — cost is a consideration but not the conceptual reason AI complements rather than replaces analysts. C is wrong — while some jurisdictions have human-in-the-loop requirements

for specific decisions, this is not the primary reason. D is wrong — modern AI security systems can and do take automated actions within defined parameters without continuous human supervision.

Q19. Which of the following is an example of adversarial AI use by a defender?

- A) Using AI to generate phishing emails to test employee security awareness
- B) Using AI to automatically approve all login requests from known IP ranges
- C) Using AI to block all encrypted traffic at the network perimeter
- D) Using AI to generate compliance reports for regulatory auditors

Answer: A **Explanation:** A is correct — simulated phishing campaigns using AI to generate realistic, personalised phishing emails are a legitimate and increasingly common defensive technique. Security teams use them to test employee awareness, measure click-through rates, and identify training gaps. This is an example of using an offensive AI capability in a controlled, authorised defensive context. B is wrong — automatically approving logins based on IP ranges is a rule-based access control decision, not adversarial AI use. C is wrong — blocking encrypted traffic is a network policy decision; this would also block legitimate traffic. D is wrong — compliance reporting is an administrative function unrelated to adversarial AI techniques.

Q20. Scenario: A security architect is presenting to the board on AI adoption. A board member asks: 'If our AI system says a login is 91% likely to be malicious, should we automatically block it?' The architect must explain the implications of acting on a probabilistic AI output. Which response BEST reflects the appropriate role of AI probabilistic outputs in security decision-making?

- A) Yes — 91% confidence is high enough that automatic blocking is always appropriate
- B) No — AI systems should never be used to make access control decisions
- C) It depends on the business context, false positive cost, and the type of action being taken — high-confidence scores can trigger automated low-impact actions, while high-impact actions should involve human review
- D) No — probabilities are unreliable and all security decisions should be based on deterministic rules

Answer: C **Explanation:** C is correct — the appropriate use of probabilistic AI outputs is context-dependent. A 91% confidence score might justify automatically stepping up authentication (requiring MFA) or adding the session to a watchlist — low-impact actions with minimal cost if the 9% uncertainty represents a legitimate user. The same score should probably not automatically block a CEO's access to a critical system mid-presentation without human review. The key variables are: what is the cost of a false positive? What is the cost of a false negative? What reversible or irreversible actions are being triggered? A is wrong — automatic blocking based on any confidence threshold will produce false positives; the 9% error rate on high-volume systems generates significant disruption. B is wrong — AI is widely and appropriately used for access control decisions in authentication and fraud detection. D is wrong — probabilistic outputs are not unreliable; they are a richer information source than binary rules.