

Statistical Learning Fundamentals

CompTIA SecAI+ | Section 2: Introduction to AI Types and Techniques

The Mathematical Foundation of AI Security

Statistical learning provides the mathematical and conceptual foundation for virtually every ML technique used in cybersecurity. Where machine learning describes the computational process of learning from data, statistical learning describes why that learning works, grounded in probability theory, estimation theory, and mathematical statistics.

Security professionals who understand statistical principles are significantly better equipped to evaluate whether a vendor's AI tool is likely to work in their environment; interpret detection confidence scores correctly; tune models to appropriate operating points; recognise when a statistical assumption has been violated causing the model to fail; and defend AI-based security decisions to auditors, regulators, and executives.

Probability: The Language of Security Risk

Every security decision is fundamentally probabilistic. Traditional security framing such as 'this port is open, therefore we are vulnerable' implies binary certainty that does not reflect reality. Statistical thinking replaces false certainty with calibrated probability estimates.

Concept	Definition	Security Application	Example
Prior Probability	Probability of an event before observing new evidence	Base rate of a threat type in your environment	P(phishing email) is approximately 0.01% of all received emails
Likelihood	Probability of observing the evidence given a hypothesis is true	How often does malware exhibit this network behaviour?	$P(\text{beacon pattern} \mid \text{malware}) = 0.85$
Posterior Probability	Updated probability after incorporating evidence via Bayes theorem	Given this alert, what is the probability of a real attack?	$P(\text{attack} \mid \text{alert}) = (0.85 \times 0.001) \text{ divided by } P(\text{alert})$
Confidence Interval	Range of values likely to contain the true parameter with stated probability	Expressing uncertainty in a risk estimate	Network anomaly score: 73 plus or minus 8 with 95% CI of 65 to 81
p-value	Probability of observing data as extreme as observed if the null hypothesis of no threat is true	Statistical significance of an anomaly	p below 0.001 means login pattern is extremely unlikely under normal behaviour

Bayes' Theorem: The Mathematics Behind Alert Confidence

Bayes' theorem is arguably the most important statistical concept for security practitioners evaluating AI tools. It calculates the posterior probability of a hypothesis given evidence: $P(H|E) = [P(E|H) \text{ times } P(H)] \text{ divided by } P(E)$. In security terms: $P(\text{attack} | \text{alert}) = [P(\text{alert} | \text{attack}) \text{ times } P(\text{attack})] \text{ divided by } P(\text{alert})$.

The counterintuitive result: even a highly accurate ML model can generate mostly false positives when the base rate of attacks is low. Example calculation: suppose a threat detection model has 99% sensitivity (recall) and 99% specificity (true negative rate). In an environment where 1 in 10,000 events is a real attack, the calculation unfolds as follows.

- For every 10,000 events: 1 true attack and 9,999 benign events
- True positives: 0.99 times 1 equals approximately 1 real attack detected
- False positives: 0.01 times 9,999 equals approximately 100 false alerts
- Positive predictive value (precision): 1 divided by (1 plus 100) equals 0.98%, meaning only 1 in 101 alerts is real

This calculation explains why SOC teams are drowning in false positives despite deploying tools marketed as 99% accurate. Bayesian thinking helps security teams correctly interpret alert confidence and design escalation workflows that account for base rates.

Normal Distribution and Baseline Establishment

The normal (Gaussian) distribution is the mathematical model underlying most anomaly detection systems in security. It is characterised by its mean (the centre of the distribution) and standard deviation (the spread). The empirical rule states: approximately 68% of observations fall within 1 standard deviation of the mean, 95% within 2 standard deviations, and 99.7% within 3 standard deviations.

In security anomaly detection, the baseline is the estimated mean and standard deviation of a metric under normal operating conditions. An observation is flagged as anomalous if it deviates beyond a defined sigma threshold. Setting that threshold involves a direct trade-off.

Threshold	False Positive Rate	Missed Detection Rate	Suitable For
Plus or minus 1 sigma	Approximately 32% of normal observations flagged	Very low; highly sensitive	Low-volume, high-consequence environments where analysts can handle high alert volume
Plus or minus 2 sigma	Approximately 5% of normal observations flagged	Low	Standard SOC operations; balanced trade-off
Plus or minus 3 sigma	Approximately 0.3% of normal observations flagged	Moderate	High-volume environments with limited analyst capacity
Plus or minus 4 sigma or more	Below 0.01% of normal observations flagged	High; misses subtle anomalies	Environments requiring near-zero false positives; risk of missing stealthy attacks

An important caveat: many real-world security metrics are not normally distributed. Network packet inter-arrival times follow heavy-tailed distributions; login frequencies may be bimodal with morning and

afternoon peaks. Applying a normal-distribution model to non-normal data produces miscalibrated anomaly thresholds. Robust anomaly detection systems test for normality and select appropriate statistical models.

Key Point: Anomaly does not equal attack. Statistical anomaly detection flags statistical outliers: events that are unusual given historical patterns. Many anomalies have legitimate explanations such as scheduled maintenance, new business initiatives, or seasonal variation. Every anomaly requires investigation; not every anomaly is a security incident.

Correlation vs. Causation in Threat Intelligence

Statistical learning excels at identifying correlations, which are statistical associations between variables. Security analysts regularly encounter correlations in threat intelligence, log analysis, and behavioural analytics. The critical skill is knowing when a correlation implies a causal mechanism that can be acted upon versus when it is spurious and acting on it wastes resources or creates blind spots.

Type	Description	Security Example	Risk of Misinterpretation
True Causal	Variable A directly causes variable B	Unpatched vulnerability causes successful exploitation	Low — action on A reduces B reliably
Confounded	Both A and B are caused by a third variable C	Ransomware attacks and HR disputes both spike during economic downturns	Medium — treating A as cause of B wastes remediation effort
Reverse Causal	B actually causes A, not vice versa	Increased IDS alerts preceded security team deploying new monitoring, not the other way around	Medium — interventions target the wrong direction
Spurious	A and B co-vary by coincidence or statistical noise	Attack frequency correlates with ice cream sales because both are seasonal	High — false hypothesis built into detection logic

In threat intelligence, the most dangerous form of misinterpreted correlation is building detection logic around indicators that are correlated with attacks in one context but not causally related to attacks in general. An attacker's C2 infrastructure may use a specific cloud provider in one campaign, but the correlation between that provider and malicious activity is real yet not causal, because millions of legitimate services use the same provider.

Regression Analysis: Quantifying Security Risk Factors

Regression analysis models the relationship between a dependent variable (what we want to predict or explain) and one or more independent variables (predictors). In security, regression is used for risk quantification and resource allocation.

Linear regression finds the equation Y equals beta-zero plus beta-one times X_1 plus beta-two times X_2 and so on, that best fits the observed data. In security, Y might be time to detect a breach and the X

variables might be number of monitored endpoints, analyst headcount, maturity of threat intelligence programme, and patch cycle frequency. The regression coefficients tell you the marginal contribution of each factor to the outcome.

Logistic regression is used when the outcome is binary such as breach versus no breach or malicious versus benign. It produces a probability between 0 and 1, making it directly applicable to classification problems. Many commercial risk scoring tools use logistic regression or regularised variants such as Lasso and Ridge that prevent overfitting as their core model, because the results are interpretable with each coefficient having a clear meaning.

Real-World Application — Capital One Breach (2019): Regression-based risk models applied to cloud misconfiguration data could have flagged Capital One's publicly accessible AWS metadata endpoint as a high-risk predictor of credential theft, a known attack vector documented in threat intelligence. The breach ultimately exposed 100 million customer records. Risk regression models that weight cloud exposure factors appropriately can surface this class of risk before exploitation.

Time Series Analysis: Detecting Temporal Attack Patterns

Time series data consists of measurements recorded at successive time intervals and is ubiquitous in security: login counts per hour, bytes transmitted per minute, CPU utilisation per second, failed authentication attempts per day. Statistical time series analysis extracts signal from noise and detects unusual temporal patterns.

- **Trend:** Long-term increase or decrease in a metric. A slowly increasing volume of outbound DNS queries over weeks may indicate progressive data exfiltration.
- **Seasonality:** Regular, predictable periodic patterns. Login volumes peak on Monday mornings and drop on weekends. A model that does not account for seasonality will generate false positives every Monday.
- **Autocorrelation:** Correlation between a time series and its own past values. Network beaconing malware produces highly regular inter-beacon intervals with extremely high autocorrelation that is statistically distinguishable from natural variation in legitimate traffic.
- **Changepoint detection:** Statistical identification of the moment a time series changes its behaviour. Changepoint algorithms can pinpoint when an attacker's reconnaissance activity transitioned to active exploitation.

Attack Pattern	Time Series Signal	Statistical Method	Detection Example
Data exfiltration	Gradual increase in outbound byte volume over days or weeks	Trend analysis, Mann-Kendall test	UEBA flags 40% week-over-week increase in outbound transfer volume
Ransomware	Sudden spike in file modification rate within seconds	Changepoint detection, control charts	EDR detects 10,000 file renames in 3 minutes, an 8 sigma deviation
C2 beaconing	Highly regular inter-beacon intervals with low variance	Autocorrelation analysis, periodicity detection	Network tool flags HTTP call every 5.000 seconds plus or minus 0.003

			seconds
Credential stuffing	Burst of failed logins across many accounts in short window	Poisson process anomaly detection	Auth log shows 5,000 failed logins across 800 accounts in 2 minutes
Slow-burn APT	Subtle deviation from normal baseline sustained over months	Long-memory time series models such as ARIMA	SIEM correlates 0.3 sigma deviation across 47 data sources over 6 months

Clustering for Threat Segmentation

Clustering algorithms group data points by statistical similarity without requiring predefined categories. In security, clustering enables the discovery of natural groupings in attack data, network traffic, and user behaviour that human analysts would not identify manually.

- K-means: Partitions data into K clusters by minimising within-cluster variance. Used for segmenting network hosts by traffic behaviour, enabling rapid identification of groups that deviate from peer behaviour. Requires pre-specifying the number of clusters.
- DBSCAN (Density-Based Spatial Clustering of Applications with Noise): Identifies arbitrarily shaped clusters and explicitly labels low-density points as noise or outliers. Particularly effective for anomaly detection because outliers are natural candidates for investigation.
- Hierarchical clustering: Builds a tree of cluster relationships. Used in malware family classification to understand evolutionary lineages between malware variants.
- Autoencoders (deep learning): Neural networks that compress data to a low-dimensional representation and reconstruct it. Reconstruction error is high for anomalous inputs, enabling anomaly detection in complex, high-dimensional data.

Feature Selection and Dimensionality Reduction

Real-world security data is high-dimensional: a single network flow record may have 50 or more attributes; a Windows event log entry may have 30 or more fields; a binary file may yield thousands of static features. Not all features are equally informative for security decisions.

Feature selection is the process of identifying which subset of available features is most predictive of the target outcome. It matters for three reasons in security contexts.

- Accuracy: Irrelevant and noisy features reduce model accuracy by introducing spurious correlations. A model trained on only the eight most predictive features of network intrusion may outperform one trained on all 50 features.
- Speed: Inference on high-dimensional data is slower. Reducing to essential features enables real-time detection on production traffic volumes.
- Interpretability: A model with eight features is far easier for an analyst to understand and audit than one with 50. Interpretability is critical for regulatory compliance and for building analyst trust in AI recommendations.

Principal Component Analysis (PCA) is the most commonly used dimensionality reduction technique. It projects high-dimensional data into a lower-dimensional space that preserves maximum variance. In security, PCA is used to visualise network traffic data where raw packet features are projected into 2D or 3D space, often revealing clear cluster separation between normal and malicious traffic that was invisible in the original high-dimensional space.

Confidence Intervals, Uncertainty, and Decision Thresholds

Statistical learning models do not produce certainties — they produce probability estimates with inherent uncertainty. Understanding how to interpret and act on uncertain predictions is a core competency for AI security practitioners.

Confidence Score	Interpretation	Recommended Action	Escalation Path
95 to 100% confidence malicious	Near-certain threat; model has high-quality signal	Automated blocking or quarantine	Log and notify analyst; no manual review required
75 to 94% confidence malicious	Strong suspicion; lower certainty	Automated isolation plus analyst investigation	Tier-1 analyst reviews within 2 hours
50 to 74% confidence malicious	Moderate suspicion; significant uncertainty	Monitoring and alerting	Tier-1 analyst reviews within 24 hours
25 to 49% confidence malicious	Weak signal; as likely benign as malicious	Log for correlation with other signals	Queue for weekly threat hunt review
Below 25% confidence malicious	Effectively noise at current signal strength	No action on single signal; aggregate over time	Include in statistical baseline monitoring

The practical implication: SOAR platforms should implement tiered response playbooks that reflect model confidence levels rather than treating all alerts as binary. High-confidence alerts trigger immediate automated responses; low-confidence alerts are batched and reviewed periodically. This approach dramatically improves SOC efficiency without increasing missed-detection risk.

Statistical Attacks on ML Models

Security professionals must also understand how adversaries exploit statistical weaknesses in ML models. Three attack classes target the statistical foundations directly.

- **Adversarial examples:** Mathematically optimised input perturbations that cause a model to misclassify with high confidence while appearing identical or near-identical to humans. In image-based security tools, adding imperceptible pixel noise can cause a malicious document to be classified as benign with 99.9% confidence.
- **Distribution shift attacks (evasion):** Attackers learn the statistical decision boundary of a deployed model and craft attack inputs that sit just outside the malicious region of the feature space, appearing statistically normal while achieving malicious objectives. Polymorphic malware that specifically generates features matching the benign distribution is an example.

- Membership inference: Given query access to a deployed model, an attacker can statistically determine whether a specific data record was in the training set. In security, this can expose that sensitive incident records or specific employee data were used to train the model, a privacy violation with regulatory implications.

Key Terms Glossary

Term	Definition
Statistical Learning	Mathematical discipline that underpins ML; uses probability and statistics to find patterns in data and quantify uncertainty in predictions
Normal Distribution	Bell-curve probability distribution characterised by mean and standard deviation; basis of many anomaly detection baselines
Standard Deviation	Measure of spread around the mean; anomaly thresholds are often expressed as multiples of standard deviations
Bayes Theorem	Formula for updating probability estimates given new evidence: $P(H E)$ equals $P(E H)$ times $P(H)$ divided by $P(E)$
Prior Probability	Probability of an event before observing evidence; the base rate of the phenomenon in the environment
Posterior Probability	Updated probability after incorporating observed evidence
Regression Analysis	Statistical method modelling the relationship between a dependent variable and independent predictors
Time Series Analysis	Statistical analysis of data measured at successive time intervals to detect trends, seasonality, and anomalies
Clustering	Unsupervised technique grouping data points by statistical similarity without predefined categories
Dimensionality Reduction	Transforming high-dimensional data to a lower-dimensional representation while preserving maximum information
PCA (Principal Component Analysis)	Linear dimensionality reduction technique projecting data onto axes of maximum variance
Adversarial Example	Input crafted with mathematically optimised perturbations to fool an ML classifier into a specific misclassification
Membership Inference	Statistical attack that determines whether a specific record was in a model's training set by analysing query responses
Changepoint Detection	Statistical method identifying the moment a time series changes its distributional properties

Quick Revision Checklist

- I can explain why statistical learning is the mathematical foundation of ML, not a separate discipline
- I can apply Bayes theorem to calculate the real-world precision of an ML security tool given its sensitivity, specificity, and base rate

- I understand the normal distribution and can select an appropriate sigma threshold given false positive and recall requirements
- I can distinguish correlation from causation in threat intelligence and describe the risk of each confusion type
- I understand how regression analysis quantifies the contribution of risk factors to security outcomes
- I can identify at least four time series patterns (trend, seasonality, autocorrelation, changepoint) with a security example for each
- I understand how clustering algorithms are used for threat segmentation and anomaly detection
- I can explain why feature selection improves model accuracy, speed, and interpretability
- I understand how to interpret model confidence scores and can describe a tiered response workflow
- I can explain adversarial examples, distribution shift attacks, and membership inference as statistical attacks on ML models
- I know why normal distribution assumptions may be violated in security data and what the consequences are

20 Exam-Based MCQs — Statistical Learning Fundamentals

Test yourself after reading. These questions are written in real exam style — scenario-heavy, with plausible distractors. Read every explanation, including for questions you answered correctly.

Q1. An ML-based threat detection model has 99% sensitivity and 99% specificity. The environment produces 1,000,000 events per day, of which 100 are real attacks. Approximately how many false positives will the model generate per day?

- A. 100
- B. 1,000
- C. 9,990
- D. 99,900

Answer: C **Explanation:** Correct: Of the 999,900 benign events, 1% (the specificity gap) will be flagged, giving approximately 9,999 false positives. With only approximately 99 true positives (0.99 times 100 attacks), the analyst faces approximately 10,098 total alerts of which 99% are false positives. This Bayesian calculation illustrates why even high-accuracy models generate analyst-overwhelming false positives at scale.

Q2. A security team observes that outbound DNS query volume has been increasing by 2% per week for the past six weeks. Which time series concept does this BEST represent, and what security hypothesis should the team investigate?

- A. Seasonality — the increase reflects normal weekly business cycles; no investigation needed
- B. Autocorrelation — past DNS queries are predicting future DNS queries; investigate DNS misconfiguration
- C. Trend — a sustained, directional increase may indicate progressive data exfiltration via DNS tunnelling

D. Changepoint — a single abrupt change has occurred; investigate for compromise on a specific date

Answer: C **Explanation:** Correct: A consistent 2% per-week increase over six weeks is a trend, which is a long-term directional movement in the time series. DNS tunnelling produces exactly this pattern as gradually increasing query volume and query length as data is encoded into DNS request fields. Seasonality (A) is periodic and predictable, not monotonically increasing. Autocorrelation (B) is about the relationship between past and future values, not a directional trend. Changepoint (D) is an abrupt break in pattern, not a gradual slope.

Q3. A threat intelligence analyst notes that 78% of observed ransomware attacks in their sector occurred between 11 PM and 3 AM on weeknights. Which statistical limitation should the analyst apply when building a time-based detection rule?

- A. This is causation — ransomware is programmed to activate at night; the detection rule is fully justified
- B. This is a correlation — the time window is associated with attacks but the mechanism (minimal staff, delayed response) should also be addressed
- C. This finding is statistically insignificant and should be disregarded
- D. Correlation implies causation when the sample size exceeds 1,000 incidents

Answer: B **Explanation:** Correct: The night-time correlation is statistically real and actionable, but the causal mechanism is reduced staff capacity to detect and respond, not something about 11 PM that causes ransomware. Simply having more analysts on night shift would address the root cause. Enhanced monitoring during those hours is justified but the underlying detection capacity must also be addressed. Option A incorrectly asserts causation. Option C is incorrect because a consistent temporal pattern is statistically meaningful. Option D is false because sample size does not convert correlation into causation.

Q4. An anomaly detection system uses a plus or minus 3 sigma threshold on network traffic volume. Which statement BEST describes the operational implication of this choice?

- A. Approximately 99.7% of normal traffic observations will be correctly cleared; the remaining 0.3% will generate false positive alerts
- B. The system will detect 99.7% of all real attacks
- C. Only traffic above 3 GB per hour will trigger alerts
- D. The 3 sigma threshold guarantees zero false positives in production

Answer: A **Explanation:** Correct: The empirical rule states that 99.7% of observations from a normal distribution fall within plus or minus 3 sigma of the mean. This means 0.3% of entirely normal traffic events will statistically fall outside the threshold and be flagged, producing false positives even with no real attacks. Option B conflates sigma thresholds with recall. Option C incorrectly equates sigma with an absolute byte threshold. Option D is false because no statistical threshold provides zero false positives.

Q5. A CISO wants to know which factors most significantly drive breach likelihood across 500 business units. Which statistical technique BEST enables this analysis?

- A. K-means clustering of business units by size
- B. Normal distribution analysis of breach frequency across the organisation
- C. Logistic regression modelling breach likelihood as a function of multiple organisational and technical risk factors
- D. Time series analysis of historical breach dates

Answer: C **Explanation:** Correct: Logistic regression is designed exactly for this problem, modelling a binary outcome such as breach or no breach as a function of multiple independent variables such as patch cadence, employee count, security training completion rate, and third-party vendor count. It produces coefficients that quantify each factor's marginal contribution to breach probability. Clustering (A) groups units but does not explain what drives breach likelihood. Normal distribution analysis (B) describes frequency but not causal factors. Time series analysis (D) is temporal, not factor-based.

Q6. An adversarial ML researcher adds imperceptible noise to a PDF file, causing a malware classifier to output 99.7% confidence benign for what is actually a malicious document. This is BEST described as:

- A. Data poisoning — malicious data was introduced during training
- B. Model extraction — the attacker is reconstructing the classifier's parameters
- C. An adversarial example — mathematically optimised input perturbation causes misclassification at inference
- D. Membership inference — the attacker is determining if the PDF was in the training set

Answer: C **Explanation:** Correct: An adversarial example is a crafted input where deliberately optimised perturbations, often imperceptible to human reviewers, cause the model to produce a high-confidence misclassification at inference time. Data poisoning (A) targets the training phase. Model extraction (B) involves reconstructing model parameters. Membership inference (D) determines whether a record was in training data.

Q7. Why does DBSCAN outperform K-means for security anomaly detection in network traffic?

- A. DBSCAN is faster to compute than K-means on large datasets
- B. DBSCAN does not require pre-specifying the number of clusters and explicitly labels low-density outlier points as anomalies
- C. DBSCAN is more accurate because it uses supervised learning
- D. DBSCAN can only be applied to network data, not user behaviour data

Answer: B **Explanation:** Correct: K-means requires pre-specifying K, the number of clusters, which is rarely known in advance for security data. DBSCAN discovers the number of clusters from the data and critically labels points that do not fit any dense cluster as noise or outliers. These outlier points are natural candidates for security investigation. DBSCAN is not necessarily faster (A) and can be slower on large datasets. It is unsupervised, not supervised (C). It is applicable to any feature-based data (D).

Q8. A logistic regression malware classifier assigns a 52% probability of malicious to a suspicious binary. Based on the tiered confidence response framework, what is the MOST appropriate SOC action?

- A. Automatically quarantine the binary and block the endpoint from the network
- B. Take no action — 52% confidence is essentially random and should be ignored
- C. Generate an alert for analyst investigation within 24 hours; do not trigger automated blocking
- D. Permanently delete the binary to eliminate any risk

Answer: C **Explanation:** Correct: At 50 to 74% confidence, the model has a moderate signal. The binary is somewhat more likely to be malicious than benign, but uncertainty is too high for automated blocking which risks disrupting legitimate systems. The appropriate response is logging an alert and queuing it for analyst review within a reasonable timeframe. Automatic quarantine (A) is appropriate only at high confidence thresholds. Ignoring 52% (B) discards genuine signal. Permanent deletion (D) without investigation is operationally destructive.

Q9. A security data scientist reduces a 50-feature network traffic dataset to 8 features using Principal Component Analysis before training an intrusion detection model. What is the PRIMARY benefit in a production environment?

- A. PCA ensures the model is 100% accurate on the training set
- B. PCA reduces model complexity, enabling faster inference on high-volume production traffic while often improving generalisation by removing noise
- C. PCA converts unsupervised models into supervised models
- D. PCA eliminates the need for labelled training data

Answer: B **Explanation:** Correct: PCA dimensionality reduction reduces computational cost at inference time, which is critical for real-time traffic analysis, and often improves generalisation by removing noisy, low-variance features that contribute spurious correlations. Option A is false because PCA does not guarantee accuracy. Option C is false because PCA is a preprocessing step that does not change the learning paradigm. Option D is false because PCA is unsupervised dimensionality reduction but the downstream classifier may still require labels.

Q10. A malware analyst discovers that a specific ransomware variant beacons home every 300 seconds with a standard deviation of 0.02 seconds. What statistical property makes this pattern detectable?

- A. High entropy — the timing is unpredictable, unlike human-generated traffic
- B. Extremely low autocorrelation — each beacon interval is independent of the previous
- C. Near-zero variance (high regularity) — machine-generated beacon intervals are statistically distinguishable from naturally variable human-initiated connections
- D. The mean interval of 300 seconds falls outside the normal distribution for network traffic

Answer: C **Explanation:** Correct: Human-initiated network connections have natural variation of seconds to minutes driven by human behaviour. A machine-generated beacon with a standard deviation of 0.02 seconds is extraordinarily regular, with a coefficient of variation far below what is achievable from

legitimate application traffic. This near-zero variance is the statistical fingerprint that network analysis tools exploit for beaconing detection. The pattern has low entropy (A), high autocorrelation (B), and the mean interval alone (D) is not the distinguishing feature.

Q11. During a threat hunt, an analyst finds that logins from a specific third-party vendor's IP range correlate with security incidents 80% of the time over the past year. What additional analysis is REQUIRED before blocking that IP range?

- A. No additional analysis — 80% correlation is sufficient justification for blocking
- B. Causal analysis: determine whether the vendor's access is the actual mechanism of compromise or whether both the vendor logins and incidents correlate with a third factor such as specific project phases
- C. Statistical analysis to confirm the correlation is not due to measurement error
- D. Legal review is required before any network block regardless of statistical evidence

Answer: B **Explanation:** Correct: An 80% correlation is operationally significant but does not establish causation. The vendor's access could be a coincidence correlated with project phases that also increase risk, or the vendor's systems could themselves be compromised, making them a victim rather than a perpetrator. Blocking the IP range based on correlation alone could disrupt legitimate business operations and miss the true root cause. Option A is incorrect because correlation alone is insufficient for attribution. Option C is a secondary concern. Option D is a compliance consideration but not the primary analytical requirement.

Q12. A statistical learning model trained to detect credential stuffing relies on Poisson distribution assumptions about login attempt rates. In practice, login attempt rates during business hours violate Poisson assumptions due to batch-scheduled authentication processes. What is the MOST likely consequence?

- A. The model will fail to train entirely
- B. The model will perform perfectly during business hours but fail at night
- C. The model will generate excess false positives during business hours when batch authentication bursts violate the Poisson rate assumption
- D. The model's precision will improve because it has more data to work with during business hours

Answer: C **Explanation:** Correct: Poisson distribution assumes events occur at a constant average rate independently of each other. Batch authentication processes produce bursts of high-rate, correlated login events that look like credential stuffing to a Poisson-based detector, generating false positives. The fix is to use a different statistical model that accounts for bursty traffic or to filter known batch-authentication sources from the model's input. Option A is false because the model trains but produces miscalibrated outputs. Option B describes the opposite pattern. Option D is false because more data does not fix a distributional mismatch.

Q13. What is the MOST significant privacy risk of training a security incident response AI on real, unredacted incident ticket data?

- A. The AI may generate incident reports that are longer than needed
- B. Membership inference attacks could allow external parties to extract specific sensitive incident details or victim identities from the trained model
- C. The model may learn to generate slower incident responses due to analysis paralysis
- D. The training data may contain duplicate tickets that reduce model accuracy

Answer: B **Explanation:** Correct: Membership inference attacks exploit the statistical properties of trained models to determine whether specific records were in the training set and in some cases to recover fragments of those records. Incident tickets containing confidential technical details, victim organisation names, or exploited vulnerability details could be partially reconstructed by anyone with query access to the deployed model. This is a real regulatory risk under GDPR and sector-specific data protection regulations. Options A, C, and D describe quality or performance issues, not privacy risks.

Q14. A security vendor's anomaly detection product displays a confidence score of 87% for a flagged login event. A junior analyst concludes there is an 87% chance this is an attack. What is wrong with this interpretation?

- A. Nothing — the analyst has correctly interpreted the confidence score
- B. The 87% figure is the model's output score, but the actual posterior probability of attack given this alert depends on the base rate of attacks in the environment, which could make the real attack probability much lower
- C. Confidence scores above 85% are always treated as confirmed attacks without investigation
- D. The analyst should only trust scores above 95% because lower scores are statistically invalid

Answer: B **Explanation:** Correct: An ML model's output score of 87% reflects the model's internal confidence based on the features of the specific event. It is not a calibrated posterior probability. The actual probability that this alert represents a real attack depends on applying Bayes theorem with the environment's base rate of attacks. In a low-attack-rate environment, even an 87% model score may correspond to only a 10 to 20% posterior probability of a real attack. This distinction between model score and posterior probability is critical for correct alert triage.

Q15. Why is feature selection important when building a network intrusion detection model with 80 available features?

- A. Using all 80 features is always optimal because more information is always better
- B. Feature selection improves model accuracy by removing noisy features, reduces inference time for real-time detection, and improves interpretability for analyst trust and regulatory review
- C. Feature selection is only required when models are deployed on embedded IoT devices with limited compute
- D. Feature selection converts supervised models into unsupervised models

Answer: B **Explanation:** Correct: Feature selection provides multiple benefits: removing irrelevant features reduces noise and overfitting which improves accuracy; reducing features reduces computational cost of inference on high-throughput network data; and producing a more interpretable model that analysts can understand and explain to auditors improves regulatory compliance. Option A is false because

irrelevant features actively degrade model performance. Option C is false because feature selection is valuable in any deployment context. Option D is false because it is a preprocessing step.

Q16. A SOAR platform is configured to automatically block source IPs when an ML model's confidence score exceeds 90%. The model occasionally outputs 92% confidence for known-good scanner IPs used by the organisation's own vulnerability assessment team. What should be done FIRST?

- A. Lower the auto-block threshold to 95%
- B. Implement an allowlist of known-good IP ranges excluded from automated blocking regardless of model confidence, contextualising the statistical model with organisational knowledge
- C. Disable the SOAR auto-block feature entirely
- D. Retrain the model to never flag vulnerability scanner IPs

Answer: B **Explanation:** Correct: The most precise and immediate fix is to allowlist known organisational scanner IPs, providing contextual knowledge that the ML model lacks. This is standard practice in security operations: ML models handle statistical pattern matching while human-curated allowlists handle known exceptions. Raising the threshold (A) reduces sensitivity across the board and may miss real threats. Disabling auto-block (C) eliminates automated response capability. Retraining (D) is the right long-term fix but takes time and may not achieve 100% exclusion.

Q17. An ML model for anomaly detection was trained assuming all features are normally distributed. Post-deployment review reveals that network connection duration follows a heavily right-skewed distribution. What is the MOST likely impact?

- A. The model's speed will degrade due to non-normal processing requirements
- B. The model's anomaly thresholds for connection duration will be miscalibrated, generating too many false positives for long-but-legitimate connections such as file transfer sessions
- C. The model will become more accurate because it has more data at the extreme end of the distribution
- D. Non-normal features do not affect ML model performance unless the model uses linear regression

Answer: B **Explanation:** Correct: Statistical models that assume normality use mean plus or minus sigma thresholds calibrated for symmetric distributions. A right-skewed distribution has a long tail of legitimate long-duration connections that fall above the 3 sigma threshold even though they are normal for that distribution. The model interprets these legitimate long transfers as anomalies, generating false positives. The fix is to apply a log-transform or use a non-parametric anomaly detection approach. Option A is false because normality assumptions do not affect processing speed. Option C is the opposite of the correct impact. Option D is false.

Q18. An attacker conducts repeated API queries against a deployed fraud detection model, testing thousands of transaction variations to find feature combinations that produce low fraud-confidence scores, then crafts actual fraudulent transactions using these feature values. This attack BEST illustrates:

- A. Data poisoning of the fraud detection model's training pipeline
- B. Distribution shift evasion — systematically probing the model's statistical decision boundary to craft inputs that fall in the benign region of the feature space
- C. Membership inference to extract training data
- D. Prompt injection via the transaction narrative field

Answer: B **Explanation:** Correct: The attacker is systematically mapping the model's decision boundary through repeated queries and then crafting transactions with feature values that the model classifies as benign. This is a distribution shift evasion attack, exploiting knowledge of the model's statistical decision surface to craft inputs that fall in the non-fraud region. Data poisoning (A) targets training. Membership inference (C) extracts training data, not the decision boundary. Prompt injection (D) is an LLM attack not applicable to tabular fraud detection models.

Q19. What does a confidence interval of '95% CI: 12 to 34 alerts per day' mean in the context of a security operations baseline?

- A. The system will alert between 12 and 34 times per day with absolute certainty
- B. If the true alert rate matches the baseline, we expect the daily alert count to fall between 12 and 34 on 95% of days; counts outside this range warrant investigation
- C. 95% of alerts in this range will be true positives
- D. The model is 95% accurate when the alert count is between 12 and 34

Answer: B **Explanation:** Correct: A 95% confidence interval for a baseline metric means that under normal operating conditions, the observed value will fall within that range 95% of the time. Days when alert counts fall outside the interval are statistically unusual and warrant investigation because either the threat landscape has changed or an operational event is affecting alert generation. Option A is incorrect because confidence intervals express probability, not certainty. Options C and D misinterpret the interval as a precision or accuracy metric.

Q20. A junior analyst reports that their ML anomaly detection model generates different numbers of alerts each Monday compared to other weekdays, causing unnecessary investigations. What statistical concept explains this pattern, and how should it be addressed?

- A. Trend — the model is gradually degrading; retrain immediately
- B. Seasonality — Monday shows higher legitimate activity that the model interprets as anomalous; the baseline should incorporate day-of-week patterns to prevent systematic false positives
- C. Autocorrelation — Friday alerts are causing Monday alerts; implement a time delay in alert generation
- D. Change point — something changed in the environment on Mondays; investigate infrastructure changes

Answer: B **Explanation:** Correct: Higher Monday activity from employees returning from weekend, scheduled batch jobs, and system sync processes is a predictable weekly seasonal pattern. A model that treats all days identically will consistently flag Monday's higher-than-average volumes as anomalous. The fix is to incorporate day-of-week seasonality into the baseline model, either by training separate baselines

per day of week or by using a seasonal decomposition model. Trend (A) describes gradual directional change, not weekly cycles. Autocorrelation (C) is a different time-series property. Change point (D) is a structural break, not a recurring pattern.