

CVE AI Working Group

CompTIA SecAI+: AI Security Certification Complete Exam Prep 2026 | Section 10: AI Threat Modeling Resources

1. What Is the CVE AI Working Group?

The Common Vulnerabilities and Exposures (CVE) program has been the cybersecurity industry's standard mechanism for tracking and communicating publicly disclosed software vulnerabilities since 1999. Every CVE identifier provides a unique, consistent label that security teams worldwide use to reference the same vulnerability—enabling coordinated patching, threat intelligence sharing, and vulnerability management at scale.

As AI frameworks and AI-enabled products have become critical business infrastructure, a practical question emerged: how should AI-related security vulnerabilities be identified, classified, and tracked within this established system? The CVE AI Working Group was formed to address precisely this question. It develops guidance for how the CVE program should handle vulnerabilities affecting AI software, AI frameworks, and AI-enabled products—bringing the same standardization discipline that has served traditional software security for decades into the AI domain.

The working group does not create a separate parallel system. Its goal is to extend the existing CVE ecosystem so that AI components can be managed through the same mature vulnerability management processes organizations already have in place for conventional software. This means AI security teams can leverage existing tooling, workflows, and institutional knowledge rather than building entirely new processes from scratch.

2. Why CVE Standardization Matters for AI Security

The value of the CVE system rests on standardization. When every stakeholder—vendors, researchers, security tools, patch management platforms, and defenders—uses the same CVE identifier to reference a specific vulnerability, the entire ecosystem becomes more efficient. A single identifier eliminates ambiguity, allows automated correlation across tools, and enables rapid communication about newly disclosed issues.

As AI frameworks such as TensorFlow, PyTorch, Hugging Face Transformers, and model serving platforms enter the CVE ecosystem, organizations gain several specific advantages.

- **Consistent identification:** the same vulnerability in an AI framework is referenced identically by vendors, researchers, and security tools—no more ambiguity about whether two reports describe the same issue.
- **Operational efficiency:** existing vulnerability management processes, ticketing systems, and patch management workflows can accommodate AI vulnerabilities with minimal modification.
- **Automated detection:** vulnerability scanners and software composition analysis (SCA) tools can detect affected AI components once they are registered in CVE databases.
- **Threat intelligence integration:** CVE identifiers allow AI vulnerabilities to appear in threat intelligence feeds, SIEM alert rules, and security dashboards alongside traditional software vulnerabilities.

- **Regulatory traceability:** compliance frameworks that require known vulnerability remediation within defined timeframes can apply the same standards to AI components when they carry CVE identifiers.

3. What Qualifies as an AI CVE: A Critical Boundary

One of the most important—and frequently misunderstood—aspects of the CVE AI Working Group's work is clarifying which AI security issues qualify for CVE assignment. Not every AI security concern represents a CVE-worthy software vulnerability.

Category	CVE Eligible?	Explanation and Examples
Exploitable implementation flaws in AI frameworks	Yes	Buffer overflows, authentication bypasses, code execution flaws in TensorFlow, PyTorch, model serving APIs
Denial-of-service vulnerabilities in AI infrastructure	Yes	Resource exhaustion flaws triggered by specially crafted model inputs to serving platforms
Authentication/authorization flaws in AI APIs	Yes	Missing authentication on inference endpoints, improper access control in MLflow or similar platforms
Model hallucinations	No	An inherent characteristic of generative AI, not an exploitable software defect in a specific product
General adversarial example susceptibility	No	A property of machine learning architectures broadly, not a specific exploitable flaw in a named product
Model bias and fairness issues	No	Ethical and design concerns, not exploitable software vulnerabilities meeting CVE criteria

This boundary matters operationally. Security teams that attempt to track AI behavioral characteristics—hallucinations, bias, general adversarial susceptibility—through the CVE system will find that these entries do not exist and will not be created. Those concerns belong in risk management frameworks like the MIT AI Risk Repository or internal AI governance processes, not the CVE vulnerability catalog.

4. The CVE Assignment Process for AI Vulnerabilities

When an AI-related vulnerability is discovered, the assignment process follows the same structure used for traditional software, with some additional judgment required at the classification stage.

- **Discovery and validation:** a researcher, security team, or vendor identifies a potential security vulnerability in an AI framework, platform, or AI-enabled product and validates that it is reproducible and meets the exploitability criteria for CVE consideration.
- **Coordinated disclosure:** the discoverer contacts the affected vendor or a CVE Numbering Authority (CNA) privately, providing sufficient technical details for the issue to be reproduced and investigated. Public disclosure is withheld until a fix is available.
- **CNA evaluation:** the CNA determines whether the reported issue meets CVE criteria—is it an identifiable, exploitable security flaw in a specific product or component? This is where the CVE AI Working Group's guidance on AI-specific boundaries is applied.

- **CVE assignment and publication:** if accepted, the vulnerability receives a unique CVE identifier. A public description is published including affected product versions, the security impact, and remediation guidance.
- **Remediation and closure:** the vendor releases a patch or mitigation. Organizations consuming the affected AI component apply the fix and document closure in their vulnerability management system.

5. Real-World Incident: AI Framework Vulnerabilities

Real-World Incident — TensorFlow Deserialization Vulnerabilities (multiple CVEs, 2021-2023): Google's TensorFlow machine learning framework received multiple CVE assignments for vulnerabilities related to unsafe deserialization of model files and specially crafted input tensors. Several of these vulnerabilities allowed denial-of-service conditions or memory corruption when processing untrusted model artifacts. These CVEs demonstrated that AI frameworks face many of the same implementation-level vulnerability classes as traditional software: input validation failures, memory safety issues, and unsafe deserialization. They also illustrated the practical value of the CVE system for AI: once assigned, these identifiers appeared in vulnerability scanners, patch management systems, and security advisories, enabling organizations to identify and remediate affected versions using existing processes.

This incident is instructive because it challenges the assumption that AI security is entirely a novel domain requiring entirely new processes. Many AI framework vulnerabilities are conventional software security flaws that happen to appear in AI-specific code. The CVE system handles these exactly as it handles vulnerabilities in any other software component.

6. Finding and Monitoring AI CVEs

Once AI vulnerabilities receive CVE identifiers, they are accessible through the same channels used to track traditional software vulnerabilities. Security teams do not need separate tools or processes to monitor AI CVEs.

- **National Vulnerability Database (NVD):** the authoritative public repository for CVE details, severity scores (CVSS), and affected product information. Search by framework name (TensorFlow, PyTorch, ONNX Runtime) or product name to identify relevant entries.
- **Vendor security advisories:** major AI framework vendors publish security advisories when new vulnerabilities are disclosed. Subscribe to advisories from vendors of frameworks, platforms, and cloud AI services your organization uses.
- **Software Composition Analysis (SCA) tools:** SCA tools scan your dependency manifests and container images for components with known CVEs, including AI frameworks and libraries. Integrate these scans into CI/CD pipelines.
- **Threat intelligence feeds:** commercial and open-source threat intelligence platforms include CVE data. Configure alerts for AI framework CVEs with high or critical severity ratings.

The key operational principle is inclusion: AI frameworks and libraries should be in your software asset inventory and should be scanned with the same frequency and rigor as web frameworks, database drivers, and operating system components.

7. Integrating AI CVEs into Vulnerability Management

The most common mistake organizations make with AI vulnerability management is treating it as a separate program. AI components belong inside the existing vulnerability management lifecycle, not alongside it.

VulnMgmt Phase	How AI CVEs Fit	Practical Action
Asset Inventory	AI frameworks are software assets	Add ML frameworks, model servers, and AI APIs to CMDB and dependency manifests
Vulnerability Scanning	SCA tools detect CVEs in AI libraries	Include AI dependencies in routine scan scope; scan container images used for AI workloads
Risk Assessment	CVSS scores apply to AI CVEs	Assess likelihood and business impact using same risk criteria as other CVEs
Prioritization	Severity + exploitability + asset criticality	Prioritize AI CVEs in security AI systems higher due to compounded risk if compromised
Remediation	Update AI framework versions	Test updated versions for model compatibility before deploying to production AI systems
Documentation	Evidence for audits and compliance	Log discovery date, assessment outcome, remediation action, and closure date

One AI-specific consideration during remediation is compatibility testing. Unlike updating a web framework where behavior changes are usually well-defined, updating an AI framework may affect model loading, inference performance, or numerical precision. Remediation planning should include a model validation step before promoting the patched version to production.

8. Responsible Disclosure for AI Vulnerabilities

The security community's responsible disclosure norms apply fully to AI vulnerabilities. Researchers and practitioners who discover exploitable flaws in AI frameworks or AI-enabled products have an ethical obligation to disclose in a way that minimizes harm to organizations that rely on those products.

Responsible Disclosure Principle: Report privately to the vendor first. Give them reasonable time—typically 90 days is the industry standard—to develop and release a fix. Coordinate on disclosure timing. Publish technical details after a patch is available so defenders can act on the information before attackers exploit it at scale.

For AI researchers this principle sometimes creates tension. Publishing novel attack techniques against AI systems (adversarial examples, new poisoning methods) is valuable academic work. The responsible disclosure framework helps navigate this: coordinate with affected framework vendors before public presentation, and where possible, publish alongside defensive guidance. The goal is to improve collective security, not to hand attackers a ready-made exploit before defenders have a chance to respond.

9. Analogy: AI CVEs and Standard Library Updates

A useful mental model for AI CVE management is the treatment of standard software library updates. When a security vulnerability is found in a widely used library—say, a JSON parsing library or an HTTP client—organizations do not create a new program for managing it. They update the library version, re-test the application, and close the vulnerability finding in their existing system.

AI framework vulnerabilities work exactly the same way. PyTorch with a known CVE is treated identically to OpenSSL with a known CVE: identify affected versions, assess exploitability in context, apply the vendor's patch, validate that the update does not break dependent applications (in this case, AI models), and document closure. The AI-specific step is the model validation—but that is an extension of existing practice, not a wholly different program.

Organizations that build this mental model early will find AI vulnerability management significantly less complex than those who approach it as an entirely separate discipline. Most of the work is already done through existing processes; the AI-specific adjustments are incremental.

10. Challenges in Classifying AI Vulnerabilities

Despite the clarity the CVE AI Working Group provides, applying CVE criteria to AI technologies introduces genuine challenges that do not arise in traditional software vulnerability assessment.

- **Scope ambiguity:** determining whether a security issue is an implementation flaw (CVE-eligible) or an inherent characteristic of ML (not CVE-eligible) requires technical judgment that may not be straightforward for novel AI attack patterns.
- **Attack vector description:** standard CVE fields (attack vector, privileges required, user interaction) were designed for traditional software. Some AI-specific attacks—such as training data poisoning via a shared dataset—do not map cleanly to these fields.
- **Remediation complexity:** for many AI vulnerabilities, the remediation is a framework version update. But some vulnerabilities may require retraining models, changing training data governance, or architectural changes—remediation paths that traditional CVE processes do not account for.
- **Rapid evolution:** new AI attack classes emerge faster than CVE processes were designed to accommodate. Emerging attack techniques may lack established CVE precedent, creating uncertainty about classification.

These challenges are being actively addressed by the working group, but practitioners should be aware that AI CVE classification involves more judgment than traditional software vulnerability assessment. When in doubt, engaging with the relevant CNA or the working group directly is appropriate.

11. Compliance Implications of AI CVEs

As regulatory frameworks increasingly address AI security, the management of AI CVEs is becoming a compliance concern alongside a security one. Organizations subject to frameworks such as the NIS2 Directive, the EU AI Act, HIPAA, or sector-specific financial regulations need to demonstrate that they identify, assess, and remediate known AI vulnerabilities with the same rigor applied to conventional software.

Practically, this means AI components must appear in the organization's software asset inventory, vulnerability scan scope, and remediation tracking systems. When auditors review vulnerability

management programs, AI frameworks should be present alongside web servers, databases, and operating systems—not absent from the evidence trail.

Organizations that establish mature AI vulnerability management processes now—before specific regulatory mandates require it—will be better positioned as AI-specific compliance requirements continue to develop. The audit evidence requirements are identical to those for traditional software: discovery date, risk assessment, remediation action, and closure date, all traceable to specific CVE identifiers.

Key Terms Glossary

Term	Definition
CVE (Common Vulnerabilities and Exposures)	A program that assigns unique standardized identifiers to publicly disclosed security vulnerabilities, enabling consistent reference across the global security ecosystem.
CVE AI Working Group	A group developing guidance for how AI-related vulnerabilities in software products and components should be identified, classified, and documented within the CVE program.
CVE Numbering Authority (CNA)	An organization authorized to assign CVE identifiers to vulnerabilities within their scope; includes major vendors, research organizations, and coordinating bodies.
National Vulnerability Database (NVD)	The US government repository that augments CVE data with severity scores (CVSS), affected product information, and remediation guidance.
CVSS (Common Vulnerability Scoring System)	A standardized framework for rating the severity of security vulnerabilities, used by the NVD and referenced in AI CVE entries as in traditional software CVEs.
Software Composition Analysis (SCA)	Automated tooling that scans software dependencies and container images to identify components with known CVEs, applicable to AI framework dependencies.
Coordinated Disclosure	The practice of privately notifying a vendor of a discovered vulnerability before public release, giving time for a patch to be developed before attackers can exploit the issue at scale.
Responsible Disclosure	An ethical framework for vulnerability reporting that prioritizes minimizing harm to users of affected products by coordinating with vendors before publishing technical details.
Implementation Flaw	An exploitable security defect introduced through specific coding or configuration decisions, as distinguished from inherent characteristics of an AI system's design.
Model Validation	Testing that confirms a trained AI model continues to behave correctly after an underlying framework is updated; an AI-specific step in the vulnerability remediation workflow.

Quick Revision Checklist

- The CVE AI Working Group extends the existing CVE program to cover AI framework and AI-enabled product vulnerabilities.
- Not all AI security concerns qualify for CVEs: implementation flaws in AI software qualify; hallucinations, bias, and general adversarial susceptibility do not.
- The CVE assignment process for AI mirrors the traditional process: discovery, coordinated disclosure, CNA evaluation, assignment, and remediation.
- Track AI CVEs using the NVD, vendor security advisories, and SCA tools integrated into CI/CD pipelines.
- AI components (ML frameworks, model servers, AI APIs) belong in the software asset inventory and routine vulnerability scan scope.
- Remediation of AI CVEs typically involves framework version updates, but must include model validation before production deployment.
- Apply the same risk assessment criteria (severity, exploitability, asset criticality) to AI CVEs as to traditional software CVEs.
- Responsible disclosure applies to AI vulnerabilities: report privately to vendors first, coordinate on timing before publishing.
- AI CVE compliance requirements are growing; maintaining audit-ready evidence of discovery, assessment, and remediation is essential.
- AI CVE classification involves more judgment than traditional software due to scope ambiguity, novel attack vectors, and rapid technology evolution.

10 Exam-Based MCQs — CVE AI Working Group

Test yourself after reading. These questions are written in real exam style — scenario-heavy, with plausible distractors. Read every explanation, including for questions you answered correctly.

Q1. Scenario: Priya is a security analyst at a financial services company responsible for vulnerability management. She receives an automated alert from the NVD indicating that a high-severity CVE has been published affecting TensorFlow 2.11, the version currently running the organization's AI fraud detection model. Which of the following actions should the security team take FIRST upon discovering that a CVE has been published for the version of TensorFlow used in the organization's AI fraud detection system?

- A) Immediately patch TensorFlow to the latest version and restart the fraud detection service
- B) Determine whether the affected version is in use, assess the business impact, and prioritize remediation
- C) Report the CVE to the CVE AI Working Group for reclassification
- D) Disable the fraud detection system until the vulnerability is resolved

Answer: B **Explanation:** B is correct because the first step in any vulnerability management process is to determine whether the affected component is in use and assess the potential business impact before taking action. A is wrong because patching without impact assessment could cause model compatibility issues or service disruption without first confirming the vulnerability is exploitable in context. C is wrong because CVEs are published after CNA review; the working group does not reclassify published CVEs through user reports. D is wrong because disabling the system before assessing whether the vulnerability is exploitable in the organization's environment is disproportionate.

Q2. Scenario: A security analyst is reviewing the organization's AI risk management program and needs to determine which concerns should be routed to the CVE-based vulnerability management system versus other governance processes. Which of the following AI security concerns would be MOST appropriate to track using the CVE system?

- A) A tendency for the organization's LLM to generate hallucinated legal citations
- B) A general susceptibility to adversarial examples in the organization's computer vision model
- C) A memory corruption vulnerability in PyTorch triggered by specially crafted tensor inputs
- D) Bias in the organization's hiring algorithm that disadvantages certain demographic groups

Answer: C **Explanation:** C is correct because a memory corruption vulnerability in PyTorch is an exploitable implementation flaw in a specific software product, which meets CVE criteria. A is wrong because hallucinations are an inherent characteristic of generative AI models, not an exploitable software defect in a specific product. B is wrong because general adversarial example susceptibility is a property of machine learning architectures broadly, not a specific named product flaw. D is wrong because algorithmic bias is an ethical and design concern tracked through AI governance processes, not the CVE system.

Q3. Scenario: A CISO is deciding how to structure the organization's response to the growing number of CVEs being assigned to AI frameworks used across the company's product portfolio. What is the PRIMARY advantage of including AI framework CVEs in an organization's existing vulnerability management program rather than managing them through a separate AI security program?

- A) AI CVEs have lower severity scores than traditional software CVEs and require less urgent attention
- B) Existing vulnerability management tools, workflows, and institutional knowledge can be applied without building entirely new processes
- C) The CVE system provides AI-specific remediation guidance not available through vendor advisories
- D) Separate AI security programs are prohibited by emerging regulatory frameworks

Answer: B **Explanation:** B is correct because integrating AI CVEs into existing vulnerability management programs allows organizations to leverage mature processes, tooling, and institutional expertise rather than building a parallel program from scratch. A is wrong because AI CVEs can carry high or critical severity ratings, and severity is independent of whether they are managed separately or within an existing program. C is wrong because CVE entries provide standardized vulnerability identification but do not offer AI-specific remediation guidance beyond what vendors provide in their advisories. D is wrong because no current regulatory framework prohibits separate AI security programs; integration is a best practice for efficiency, not a legal requirement.

Q4. Which of the following BEST describes the role of a CVE Numbering Authority (CNA) in the AI vulnerability disclosure process?

- A) The CNA develops patches for AI vulnerabilities and distributes them to affected organizations
- B) The CNA evaluates reported vulnerabilities, determines CVE eligibility, and assigns CVE identifiers

- C) The CNA publishes public proof-of-concept exploit code to accelerate vendor response
- D) The CNA is responsible for classifying all AI risks, including hallucinations and model bias

Answer: B **Explanation:** B is correct because CNAs are authorized organizations that evaluate reported vulnerabilities, apply CVE criteria to determine eligibility, and assign CVE identifiers to qualifying issues. A is wrong because CNAs do not develop patches; remediation is the responsibility of the affected vendor. C is wrong because CNAs do not publish exploit code; they publish vulnerability descriptions that support coordinated remediation. D is wrong because CNAs evaluate software implementation flaws against CVE criteria; inherent AI characteristics like hallucinations and bias are outside CVE scope.

Q5. An organization's SCA scanning tool reports no CVEs for the AI frameworks in its environment. What is the MOST important caveat to communicate about this finding?

- A) The organization's AI systems are fully secure and require no further vulnerability assessment
- B) SCA tools only detect CVEs for registered components; vulnerabilities in AI models themselves or in emerging attack classes may not have CVE entries yet
- C) The SCA tool is malfunctioning and needs to be replaced with an AI-specific vulnerability scanner
- D) CVEs are not applicable to AI systems, confirming that traditional vulnerability management does not need to cover AI

Answer: B **Explanation:** B is correct because a clean SCA scan means no known CVEs were detected in the scanned components, but this does not eliminate AI-specific risks that fall outside CVE scope (hallucinations, adversarial susceptibility, data poisoning risks) or newly discovered vulnerabilities not yet assigned a CVE. A is wrong because no CVE finding is a point-in-time result for software vulnerabilities, not a comprehensive security assessment. C is wrong because SCA tools are appropriate for detecting AI framework CVEs; the caveat is about scope, not tool malfunction. D is wrong because CVEs are applicable to AI frameworks, and traditional vulnerability management must extend to include AI components.

Q6. A security researcher discovers an exploitable authentication bypass in a commercial AI model serving platform. What is the MOST appropriate next step under responsible disclosure principles?

- A) Publish full technical details immediately to maximize community awareness
- B) Report the vulnerability privately to the vendor and allow time for a patch before public disclosure
- C) Submit the vulnerability directly to the NVD for immediate publication
- D) Demonstrate the exploit at an upcoming security conference without notifying the vendor first

Answer: B **Explanation:** B is correct because responsible disclosure requires private notification to the vendor before public release, giving them time—typically 90 days—to develop and release a fix so defenders can protect their systems before attackers can exploit widely published technical details. A is wrong because immediate public disclosure before a patch is available gives attackers an advantage over defenders. C is wrong because the NVD does not accept direct vulnerability submissions; they augment CVE data after assignment by CNAs. D is wrong because conference disclosure without vendor notification violates coordinated disclosure norms and is considered unethical in the security community.

Q7. Why must organizations include a model validation step when remediating AI framework CVEs, unlike typical software library updates?

- A) Because AI framework updates always break model compatibility and require full model retraining
- B) Because updated AI frameworks may affect model loading, inference behavior, or numerical precision that must be verified before production deployment
- C) Because CVE remediation for AI frameworks requires regulatory approval before implementation
- D) Because AI frameworks cannot be patched; the entire model must be replaced when a CVE is disclosed

Answer: B **Explanation:** B is correct because AI framework version changes can affect model loading compatibility, numerical computation results, or inference performance in ways that may not be immediately obvious, requiring validation testing before the update is promoted to production. A is wrong because updates do not always break compatibility or require retraining; the model validation step determines whether they do. C is wrong because CVE remediation for AI frameworks does not require regulatory approval before patching. D is wrong because AI framework CVEs are typically addressed through framework version updates, not model replacement.

Q8. Which of the following represents the PRIMARY challenge the CVE AI Working Group addresses regarding AI vulnerability classification?

- A) Determining which CVEs apply to traditional software versus AI software
- B) Clarifying the boundary between exploitable AI software implementation flaws and inherent AI characteristics
- C) Assigning CVE identifiers to model hallucinations and bias issues
- D) Replacing the OWASP Top 10 with a CVE-based risk prioritization framework

Answer: B **Explanation:** B is correct because the working group's primary contribution is providing guidance on where the line falls between CVE-eligible implementation flaws (exploitable software defects in specific products) and inherent AI characteristics (hallucinations, general adversarial susceptibility, bias) that do not meet CVE criteria. A is wrong because the challenge is not distinguishing AI from traditional software CVEs but rather determining what AI security issues qualify as CVEs at all. C is wrong because hallucinations and bias do not meet CVE criteria and the working group does not assign them CVE identifiers. D is wrong because the working group's scope is CVE classification guidance, not developing a competing risk prioritization framework.

Q9. An organization is preparing for a compliance audit under a framework that requires evidence of known vulnerability remediation. Which of the following BEST demonstrates mature AI CVE management?

- A) A list of all AI risks identified in the MIT AI Risk Repository
- B) Documentation showing AI framework CVE discovery dates, risk assessments, remediation actions, and closure with traceable CVE identifiers
- C) Confirmation that the organization's AI models have never experienced adversarial attacks

D) A vendor statement that all AI frameworks used are free of known vulnerabilities

Answer: B **Explanation:** B is correct because audit evidence for vulnerability management requires a documented trail: the vulnerability was discovered, risk was assessed, remediation was performed, and the finding was closed—all traceable to specific CVE identifiers. This is the standard evidence package for both traditional and AI CVE compliance. A is wrong because the MIT AI Risk Repository identifies risk categories, not specific CVE-tracked software vulnerabilities. C is wrong because absence of known attacks does not demonstrate that vulnerabilities were managed; attackers may simply not have targeted the system yet. D is wrong because vendor statements are not audit evidence of an organization's own remediation activities.

Q10. Which of the following BEST completes the picture of AI security when the CVE ecosystem is combined with OWASP frameworks and the MIT AI Risk Repository?

- A) The CVE ecosystem eliminates the need for OWASP frameworks by providing a more precise vulnerability catalog
- B) The CVE ecosystem provides standardized tracking and remediation of AI software implementation flaws, complementing risk frameworks that address broader AI security concerns
- C) The CVE ecosystem provides attack simulation capabilities that replace MITRE ATLAS red team exercises
- D) The CVE ecosystem covers all AI security concerns, making the MIT AI Risk Repository redundant

Answer: B **Explanation:** B is correct because the CVE ecosystem handles a specific and well-defined scope: identifiable, exploitable implementation flaws in AI software products. This complements OWASP frameworks (which prioritize critical risk categories) and the MIT repository (which catalogs broad AI risks including those outside CVE scope). Each resource addresses a different layer of the AI security landscape. A is wrong because OWASP frameworks cover risk prioritization and security guidance beyond what the CVE catalog provides. C is wrong because CVE tracking is a vulnerability management process, not an attack simulation framework; MITRE ATLAS serves a fundamentally different purpose. D is wrong because the MIT repository covers risks—including inherent AI characteristics—that the CVE system explicitly does not cover.