

AI Governance Engineer

1. From Policy Statements to Operational Reality

Organizations can author comprehensive AI governance policies — defining acceptable use cases, documentation requirements, human oversight mandates, and bias evaluation thresholds. Those policies achieve nothing if no technical mechanism enforces them. The AI Governance Engineer bridges the gap between policy intent and operational reality by translating governance requirements into system configurations, workflow controls, automated compliance checks, and audit mechanisms that make policy violations detectable and compliance measurable.

The role has grown urgent as AI regulation has moved from aspirational to legally binding. The EU AI Act (2024), NIST AI RMF 1.0, and ISO/IEC 42001:2023 each impose specific compliance obligations that require technical implementation, not just documentation. Organizations without AI Governance Engineers discover this the hard way during regulatory examination, when they find their policies are aspirational rather than enforced. This is not a compliance checkbox role — it is an engineering role that produces infrastructure other functions depend on.

2. The Regulatory Landscape: Three Frameworks in Scope

Framework	Type	Jurisdiction	Key Requirement for Governance Engineers
EU AI Act (2024)	Binding regulation	EU — extraterritorial	Conformity assessment for high-risk AI; prohibited-use enforcement; post-market monitoring
NIST AI RMF 1.0	Voluntary framework	US — widely adopted	GOVERN, MAP, MEASURE, MANAGE functions; documentation and accountability requirements
ISO/IEC 42001:2023	Management system standard	International — certifiable	AI Management System implementation; internal audit; continual improvement
GDPR	Binding regulation	EU — extraterritorial	Data Protection Impact Assessments (DPIAs) for AI processing personal data at scale

AI Governance Engineers operating in multinational organizations must simultaneously satisfy overlapping frameworks. A model deployed in healthcare in the EU may need to address EU AI Act Annex III, GDPR Article 35 (DPIA), and ISO/IEC 42001 management system requirements — three audit regimes with partially overlapping documentation standards. The governance engineer's job is to design a single documentation and control architecture that satisfies all three simultaneously, avoiding duplicated effort while leaving no compliance gap.

3. EU AI Act Risk Classification: Practical Implications

The EU AI Act risk-based classification system is the most consequential regulatory framework AI Governance Engineers implement. Article 5 prohibits specific AI applications outright — social scoring by public authorities, real-time biometric identification in public spaces (with narrow exceptions), and AI systems that exploit psychological vulnerabilities. Annex III lists high-risk applications requiring full conformity assessment before deployment: AI in critical infrastructure, employment decision-making (CV screening, interview assessment), access to essential services (credit scoring, insurance), and law enforcement applications.

- **Prohibited (Article 5):** No deployment path exists. AI Governance Engineers implement hard deployment blocks — not policy guidance — for prohibited-category applications. This means the model registry must include a risk classification field and the deployment gate must refuse promotion for prohibited-category models regardless of business pressure.
- **High-risk (Annex III):** Mandatory conformity assessment, EU database registration, technical documentation, human oversight mechanisms, and post-market monitoring. This is the compliance-heavy category that most governance engineering effort addresses.
- **Limited-risk:** Transparency obligations only — chatbots must identify themselves as AI; deepfakes must be labeled as synthetic. Lower technical implementation burden but still requires disclosure controls.
- **Minimal-risk:** No specific obligations under the EU AI Act. Most consumer-facing and internal AI applications fall here. Governance programs typically still apply internal standards even where regulation does not mandate them.

4. Technical Implementation of Governance Controls

Governance engineering is infrastructure work, not policy writing. AI Governance Engineers implement controls across the AI development and deployment stack. The following describes the control layers a mature AI governance infrastructure includes:

- **Model registration workflows:** Governance metadata — intended use case, risk classification, training data documentation, fairness evaluation results — is required before a model can enter the registry. The registration system rejects submissions with missing required fields automatically, without human intervention.
- **Deployment gates:** CI/CD pipeline checks verify that governance requirements are satisfied before any model is promoted to a higher environment. Fairness metrics below threshold, missing model cards, or incomplete regulatory documentation cause the pipeline to fail and block promotion.
- **Audit logging:** Governance-relevant events are captured in an immutable log — who approved a model for production, what documentation was submitted, when governance reviews occurred, which version was promoted when. This logging must be tamper-evident to serve as valid audit evidence.
- **Compliance dashboards:** Governance stakeholders have real-time visibility into the AI portfolio — deployed models, risk categories, review status, monitoring health — enabling proactive identification of models approaching review deadlines or drifting from compliance thresholds.
- **Automated fairness monitoring:** Post-deployment fairness metrics are monitored continuously against production inference data. Alert thresholds trigger human review when demographic performance disparities exceed the tolerances documented in the model's governance artifacts.

Design Principle: Technical enforcement beats policy compliance. Governance controls that require practitioners to voluntarily follow policy are always weaker than controls that make non-compliance technically impossible or automatically detectable. A deployment gate that cannot be bypassed is worth more than ten policy reminders.

5. Model Cards and Data Cards as Governance Artifacts

Model cards and data cards are the central documentation artifacts that AI Governance Engineers implement, maintain, and audit. Google Research introduced the model card format in 2019; it has since been codified into NIST AI RMF documentation requirements and referenced in EU AI Act technical documentation requirements (Article 11). Organizations that implement model cards early find that they become a forcing function for good engineering practice — teams must explicitly document limitations and fairness evaluations that might otherwise be glossed over.

Artifact	Core Contents	When Required	Standard Tooling
Model Card	Intended use, performance metrics, fairness evaluation, limitations, training approach, version history	Before deployment; updated at each model version	Hugging Face model cards, Google Model Card Toolkit, Microsoft Responsible AI Toolbox
Data Card	Dataset sources, collection method, preprocessing, known biases, data governance status, retention policy	Before model training; updated if dataset changes	Google Data Cards Playbook, Croissant metadata format
System Card	End-to-end description for multi-model pipelines; failure mode analysis; human oversight mechanisms	Complex AI systems; EU AI Act Article 11 technical documentation	Organization-specific templates

AI Governance Engineers implement tooling that generates draft model and data cards from system metadata automatically, then requires human review and formal sign-off before the artifact is accepted as the governance record. This reduces documentation friction while maintaining the human accountability that governance requires. The combination of automation for draft generation and human sign-off for acceptance is the design pattern that satisfies both efficiency and compliance requirements.

6. Bias Detection and Fairness Monitoring Tools

Tool	Provider	Primary Function	Integration Context
AI Fairness 360 (AIF360)	IBM Research	Comprehensive fairness metrics; bias mitigation algorithms	Python pipelines; pre-deployment evaluation
Fairlearn	Microsoft	Fairness evaluation; constrained optimization	Azure ML; scikit-learn compatible
What-If Tool	Google	Interactive model behavior exploration across groups	TensorFlow; auditor review workflows

SageMaker Clarify	Amazon AWS	Bias detection in data and outputs; explainability reports	AWS SageMaker pipelines; automated EU AI Act documentation
Aequitas	Univ. of Chicago	Algorithmic discrimination audit; disparity metrics	Standalone Python; public-sector audit programs

Post-deployment fairness monitoring is required under the EU AI Act for high-risk applications — a single pre-deployment evaluation is legally insufficient. AI Governance Engineers configure monitoring infrastructure that re-evaluates fairness metrics on production inference data at defined intervals, comparing real-world demographic performance against the thresholds documented in the model's approved governance artifacts. When metrics degrade, the alert must reach a human reviewer who can investigate before the system causes material harm.

7. Shadow AI: The Governance Gap That Technical Controls Close

Organizational Pattern — Shadow AI at Scale (2024): Gartner research in 2024 found that more than 40% of AI projects at large enterprises were initiated outside formal IT and governance channels. Teams used commercial AI APIs directly in production workflows without approved use-case documentation, risk assessment, data handling review, or vendor security evaluation. The organization had no visibility into what sensitive customer or employee data was being transmitted to external AI providers. This is not a sophisticated attack; it is an organizational gap that creates exactly the exposure an adversary would exploit and a regulator would penalize. AI Governance Engineers close this gap through technical controls, not policy reminders.

The technical response to shadow AI requires controls at three levels:

- **Network controls:** All external AI API traffic must traverse an approved proxy. Connections to unapproved AI API endpoints are blocked at the network layer — not merely discouraged by policy. The proxy logs all requests, providing visibility into which AI services organizational data is being sent to.
- **Discovery tooling:** Browser extension analysis, network traffic monitoring, and endpoint scanning identify unauthorized AI tool usage. An approved-tool catalog with clear procurement paths makes compliant tool choices easier than unauthorized alternatives.
- **Enforcement with consequences:** Governance policies include defined consequences for shadow AI deployment. Detected violations and enforcement actions create deterrence that guidance documents alone cannot generate.

8. The NIST AI RMF GOVERN Function: Technical Implementation

GOVERN Category	Requirement	Technical Implementation
GOVERN 1	Policies, processes, and procedures for AI risk management established and enforced	Model registration workflows; deployment gates; policy-as-code in CI/CD pipelines

GOVERN 2	Accountability structures for AI risk defined and operational	Approval workflows with identity-linked sign-off; immutable audit logs capturing accountability events
GOVERN 3	Organizational culture and workforce competency cultivated	Training requirement gates; system access conditional on completed training records
GOVERN 4	Teams committed to AI risk management	Governance metrics dashboards visible to leadership; escalation workflows with documented SLAs
GOVERN 5	Organizational risk policies encompass AI risks	Risk register integration; governance metadata linked to risk management systems
GOVERN 6	AI risk policies reflect organizational values	Use-case classification aligned to ethics policy; hard blocks for prohibited-use categories

9. Working with Adjacent Roles: Governance as Connective Tissue

The AI Governance Engineer does not work in isolation — this role sits at the center of a network of AI risk and security functions. Understanding these relationships is essential for SecAI+ exam scenarios that describe cross-functional AI governance challenges.

- **AI Risk Analyst:** The Risk Analyst identifies and quantifies AI risks; the Governance Engineer implements controls to mitigate those risks and monitors whether controls are working. Risk assessments should directly update governance control thresholds — the two functions operate in a feedback loop, not a one-time handoff.
- **AI Security Architect:** The Security Architect designs the technical security architecture for AI systems. The Governance Engineer ensures that security controls are documented in model cards and are reflected in governance artifacts — security and governance documentation must be consistent.
- **AI Auditor:** The Auditor provides independent verification that governance controls are operating. The Governance Engineer is the primary source of audit evidence — audit logs, model registry records, and deployment gate histories all originate in governance infrastructure.
- **Data Engineer:** Data Engineers produce the training data that models depend on. Governance Engineers ensure that data provenance and governance documentation (data cards) accompany training datasets into the model development workflow.

10. Red Team Perspective: Exploiting Governance Gaps

Red teams testing AI governance programs probe systematically for gaps that technical enforcement would close. The most common governance vulnerabilities are process failures — not sophisticated technical exploits — because governance often relies on voluntary practitioner compliance rather than technical enforcement.

- Models in the registry with missing or placeholder metadata: Registration is not technically enforced — metadata can be submitted incompletely without blocking deployment.
- Shadow AI traffic detectable via network analysis: Network-level enforcement is absent — governance relies on practitioners self-reporting tool usage.
- Automated pipelines with no human approval step: Governance gates are not integrated into the deployment process — models can reach production without governance sign-off.
- Audit log gaps for governance event types: Logging was not designed to capture governance-relevant events — policy violations are invisible to retrospective review.
- Fairness results not cryptographically linked to deployed model versions: Governance artifacts can be decoupled from the models they document — compliant documentation can be attached to non-compliant models.

11. Responsible Escalation: When Governance Engineers Must Refuse

AI Governance Engineers have visibility across the organization's AI portfolio. This visibility creates obligations that cannot be discharged passively. The EU AI Act's personal liability provisions make these obligations legally enforceable, not merely professional best practice.

- Refuse and document formally: An AI Governance Engineer asked to approve a model that failed fairness evaluation must refuse and document the refusal — not grant conditional approval that circumvents the governance gate.
- Escalate policy violations: Discovering a model deployed outside the governance process requires formal escalation to governance leadership, not a note in a report that no one reads.
- Personal legal exposure: The EU AI Act creates personal liability for individuals who knowingly approve non-compliant high-risk AI systems. Governance engineers who treat approval as a rubber stamp create personal legal risk alongside organizational risk.
- Pre-defined escalation paths: Named escalation contacts and documented response SLAs make governance escalation a normal operational practice rather than a career-limiting action.

Memory Anchor — Building Inspector: An AI Governance Engineer is the building inspector of the AI production environment. A building inspector who approves a structurally deficient building because the developer applied pressure has failed professionally and legally. The governance engineer's approval gate carries the same obligation: the approval is a certification, not a formality.

Key Terms Glossary

Term	Definition
EU AI Act (2024)	EU regulation establishing risk-based AI classification; mandatory conformity assessment for high-risk applications; personal liability for knowing non-compliance approval
Conformity Assessment	Formal verification that a high-risk AI system meets EU AI Act requirements before deployment; may require independent third-party review for highest-risk applications
NIST AI RMF 1.0	US voluntary AI risk management framework organized around four

	functions: GOVERN, MAP, MEASURE, MANAGE
ISO/IEC 42001:2023	International certifiable management system standard for AI; analogous to ISO 27001 for information security; requires internal audit and continual improvement
Model Card	Standardized documentation artifact describing a model's intended use, performance metrics, fairness evaluations, and limitations
Data Card	Documentation artifact describing training dataset characteristics — sources, collection methods, known biases, preprocessing steps, and data governance status
Shadow AI	AI tools or models deployed outside formal IT and governance channels — creating undocumented risk exposure and regulatory compliance gaps
Deployment Gate	Automated CI/CD pipeline check that blocks model promotion when governance requirements (documentation, fairness thresholds, approvals) are not satisfied
AI Fairness 360 (AIF360)	IBM open-source toolkit providing comprehensive fairness metrics and bias mitigation algorithms for machine learning models
Disparate Impact	Statistical measure of whether model outcomes differ significantly across protected demographic groups — a key fairness metric in EU AI Act and US employment law
Post-Market Monitoring	EU AI Act requirement for ongoing production monitoring of high-risk AI system performance and fairness — a one-time pre-deployment evaluation is insufficient
GOVERN Function	NIST AI RMF function covering organizational policies, accountability structures, and culture for AI risk management; implemented technically by the AI Governance Engineer

Quick Revision Checklist

- AI Governance Engineers translate policy into technical controls — deployment gates, registration workflows, audit logging, compliance dashboards.
- EU AI Act: four tiers — prohibited (Article 5), high-risk (Annex III), limited-risk, minimal-risk — each with distinct governance obligations.
- Model cards and data cards are central governance artifacts; governance engineers implement tooling to generate drafts and enforce human sign-off.
- Fairness tools: AIF360 (IBM), Fairlearn (Microsoft), What-If Tool (Google), SageMaker Clarify (AWS), Aequitas (UChicago).
- Shadow AI is closed by network-layer proxy enforcement and discovery tooling — not policy reminders.
- NIST AI RMF GOVERN function has six categories (GOVERN 1-6) implemented technically by governance engineers.
- EU AI Act post-market monitoring requires ongoing production fairness evaluation — pre-deployment evaluation alone is legally insufficient for high-risk applications.
- Personal legal liability under EU AI Act for knowingly approving non-compliant high-risk AI systems.

- Technical enforcement beats policy compliance — deployment gates that cannot be bypassed outperform voluntary compliance.
- Governance Engineers work with Risk Analysts (risk identification), Security Architects (control design), Auditors (evidence provision), and Data Engineers (data provenance).
- AI Governance Engineers formally refuse and document when a model fails a required evaluation — conditional approvals circumvent the governance gate.

10 Exam-Based MCQs — AI Governance Engineer

These questions are written in SecAI+ exam style. Read every explanation, including for questions you answered correctly.

Q1. *Scenario:* A financial services firm has an AI governance policy requiring model cards, fairness evaluations, and risk classification before deployment. Six months after implementation, a compliance audit finds three production models with missing documentation. The CISO asks the AI Governance Engineer what change would prevent this in future. Which action BEST demonstrates that an AI governance program is technically enforced rather than policy-dependent?

- A) Publishing a comprehensive AI acceptable use policy on the intranet
- B) Requiring practitioners to complete annual AI governance training
- C) Configuring deployment pipelines to block model promotion when governance metadata is incomplete
- D) Conducting quarterly governance reviews with the CISO

Answer: C **Explanation:** C is correct because deployment gates make policy compliance automatic — a model without complete governance metadata cannot be deployed regardless of practitioner intent. A is wrong because publishing a policy creates awareness but no enforcement mechanism. B is wrong because training improves knowledge but cannot prevent non-compliant deployments. D is wrong because quarterly reviews are detective controls — they identify violations after deployment, not prevent them.

Q2. *Scenario:* During routine monitoring, an AI Governance Engineer finds a production model not in the model registry and with no governance documentation. Investigation reveals the business unit used direct infrastructure credentials to deploy it. The model processes customer loan applications. An AI Governance Engineer discovers that a business unit deployed a model to production by bypassing the CI/CD pipeline governance gates. What is the MOST appropriate immediate response?

- A) Document the incident in the risk register and address it in the next quarterly review
- B) Remove the model from production and escalate to governance leadership with formal documentation
- C) Retroactively complete the model card and fairness evaluation for the deployed model
- D) Issue a reminder to the business unit about governance policy requirements

Answer: B **Explanation:** B is correct because a model deployed without governance review has not been assessed for risk, fairness, or regulatory compliance — it must be removed and the bypass formally escalated. A is wrong because deferring to a quarterly review allows a potentially non-compliant model to remain in production. C is wrong because retroactive documentation does not remediate the compliance

gap — the model was deployed without the review that documentation is supposed to precede. D is wrong because a policy reminder is not commensurate with a governance bypass that may have regulatory implications under the EU AI Act.

Q3. Scenario: A healthcare AI company deploys two models: a high-risk AI system for radiology image analysis (Annex III) and a minimal-risk AI system for appointment scheduling. The AI Governance Engineer is defining ongoing compliance obligations for each. Under the EU AI Act, which activity is REQUIRED for high-risk AI applications but NOT specifically required for minimal-risk applications?

- A) Model card documentation
- B) Post-market monitoring of model performance and fairness
- C) NIST AI RMF GOVERN function implementation
- D) Acceptable use policy publication

Answer: B **Explanation:** B is correct because the EU AI Act specifically requires post-market monitoring for high-risk applications under Article 72 — ongoing production monitoring, not a one-time pre-deployment evaluation. A is wrong because model cards are a best practice for all AI systems; the EU AI Act's specific technical documentation requirements apply to high-risk applications but model card practices are not mandated for minimal-risk applications by the Act itself. C is wrong because NIST AI RMF is a US voluntary framework, not an EU regulatory requirement. D is wrong because acceptable use policy is an internal governance practice not defined by EU AI Act requirements.

Q4. Which tool is MOST appropriate for measuring whether an AI model's outputs show statistically significant disparate impact across demographic groups before deployment?

- A) MLflow Model Registry
- B) IBM AI Fairness 360 (AIF360)
- C) Hugging Face model cards
- D) AWS CloudTrail

Answer: B **Explanation:** B is correct because AIF360 provides a comprehensive library of fairness metrics designed to measure disparate impact and other bias indicators across demographic groups, plus bias mitigation algorithms that can be applied when disparities are identified. A is wrong because MLflow is a model versioning and tracking tool — it manages model artifacts but does not evaluate fairness metrics. C is wrong because model cards are documentation artifacts — they are populated with fairness results but do not generate them. D is wrong because CloudTrail is an AWS API audit logging service — it tracks governance events but does not perform statistical fairness evaluation.

Q5. An organization implements tooling that automatically generates draft model cards from model metadata. What critical step must be included to satisfy governance accountability requirements?

- A) Storing draft model cards in an encrypted repository
- B) Requiring human review and formal sign-off before the model card is accepted as the governance artifact

- C) Generating model cards in a read-only format to prevent unauthorized edits
- D) Including model card generation as a step in the data ingestion pipeline

Answer: B **Explanation:** B is correct because automated generation produces a draft from metadata — human review and formal sign-off verify accuracy, confirm that limitations are honestly stated, and create the accountability record that governance requires. The combination of automation for drafting and human acceptance for the final artifact satisfies both efficiency and compliance. A is wrong because encryption is a data protection control, not an accountability mechanism for artifact quality. C is wrong because preventing edits does not ensure the content is accurate or that accountability has been accepted. D is wrong because model card generation belongs in the model training and deployment workflow, not the data ingestion pipeline.

Q6. Which NIST AI RMF GOVERN category is MOST directly addressed by implementing identity-linked approval workflows for model deployments?

- A) GOVERN 1 — Policies and procedures established
- B) GOVERN 2 — Accountability structures operational
- C) GOVERN 3 — Workforce competency cultivated
- D) GOVERN 6 — Policies reflect organizational values

Answer: B **Explanation:** B is correct because GOVERN 2 requires accountability structures to be defined and operational — identity-linked approval workflows make accountability observable and auditable by recording which specific individual approved each governance action. A is wrong because GOVERN 1 addresses the existence and communication of policies, not the specific accountability mechanisms within them. C is wrong because GOVERN 3 addresses training and competency development — approval workflows assume competency; they do not build it. D is wrong because GOVERN 6 addresses policy alignment with organizational values — a policy design function, not an implementation mechanism.

Q7. An AI Governance Engineer is asked to approve a model that failed the organization's required fairness evaluation threshold. The business sponsor argues the model is urgently needed for a product launch. What is the CORRECT action?

- A) Approve the deployment with a condition that the fairness issue be remediated within 90 days
- B) Escalate to the CISO for a risk acceptance decision
- C) Formally refuse approval and document the refusal in the governance record
- D) Grant a temporary exception if the business impact of delay exceeds a defined threshold

Answer: C **Explanation:** C is correct because the governance gate exists precisely for this situation — a model that failed a required evaluation must not be approved, and the refusal must be formally documented to create an audit trail. Under the EU AI Act, knowingly approving a non-compliant high-risk application creates personal legal liability. A is wrong because conditional approval of a failed model circumvents the governance gate — the model is in production in a non-compliant state regardless of the remediation condition. B is wrong because governance gate failures should not be reframed as CISO risk acceptance decisions — the gate exists to prevent exactly this escalation path. D is wrong because ad hoc exceptions based on business impact are a governance failure pattern; exceptions must be pre-defined in

policy.

Q8. Shadow AI deployments create which PRIMARY governance risk for an organization?

- A) Increased model training costs from duplicated compute resources
- B) Absence of documentation, risk assessment, and data handling review for AI tools in production
- C) Performance degradation from unauthorized models competing for inference resources
- D) Exposure of organizational model intellectual property to external AI providers

Answer: B **Explanation:** B is correct because shadow AI deployments bypass the governance process — no use-case documentation, risk assessment, fairness evaluation, or data handling review has occurred, meaning the organization has no understanding of what risks these tools create or what regulatory obligations they trigger. A is wrong because compute cost duplication is an operational concern, not the primary governance risk. C is wrong because resource competition is an operational concern distinct from the governance gap created by unapproved deployments. D is wrong because IP exposure is a possible shadow AI risk, but the primary governance risk is absence of review for AI tools that may be consuming and transmitting sensitive organizational data.

Q9. Which control MOST effectively detects unauthorized external AI API usage within an enterprise network?

- A) Annual employee surveys about AI tool usage
- B) Network traffic monitoring with approved AI API proxy enforcement
- C) Quarterly reviews of software license procurement records
- D) Code repository scanning for AI library imports

Answer: B **Explanation:** B is correct because requiring external AI API traffic to traverse an approved proxy provides both detection (proxy logs all AI API calls) and enforcement (unapproved endpoints blocked at the network layer), closing the shadow AI gap at the infrastructure level. A is wrong because employee surveys are self-reporting mechanisms — practitioners using shadow AI tools are unlikely to self-report non-compliant usage. C is wrong because shadow AI often uses consumption-based APIs paid on business credit cards, invisible to software procurement records. D is wrong because code scanning detects library imports in committed code — it misses HTTP-based API usage without imported libraries and does not address tools used by non-developers through web interfaces.

Q10. The EU AI Act creates which personal consequence for individuals who knowingly approve non-compliant high-risk AI systems for deployment?

- A) Mandatory remediation training with competency re-certification
- B) Personal legal liability
- C) Suspension of AI governance professional certification
- D) Organizational audit trigger with no personal consequences for the individual

Answer: B **Explanation:** B is correct because the EU AI Act creates personal legal liability for individuals who knowingly approve non-compliant high-risk AI systems — a significant departure from regulatory

frameworks that focus exclusively on organizational penalties. This is why formal documentation of refusals and the governance gate function matters beyond organizational compliance. A is wrong because the EU AI Act creates legal consequences, not training requirements, for knowing non-compliance approval. C is wrong because AI governance certification suspension is not an EU AI Act mechanism. D is wrong because the EU AI Act does create personal consequences — it does not limit consequences to the organization.