

Accountability

CompTIA SecAI+: AI Security Certification Complete Exam Prep 2026 | Section 24: Responsible AI

1. Accountability: The Missing Link in AI Governance

When AI systems cause harm, who is responsible? This question, deceptively simple, reveals one of the deepest governance challenges in AI deployment. In 2020, a Dutch court ordered the shutdown of SyRI (System Risk Indication), a government AI that analysed multiple government datasets to identify potential welfare fraud. SyRI disproportionately flagged residents of lower-income, heavily immigrant neighbourhoods. Individuals flagged faced intrusive investigations and loss of benefits. The algorithm made the selection. The data scientists didn't investigate individuals. The agency said it was following AI recommendations. The court found privacy and human rights violations — but the accountability chain was diffuse enough that no individual or team had been formally designated responsible for the system's fairness or its outcomes.

Accountability in AI is the requirement that identifiable parties — individuals, teams, organisations — bear responsibility for AI outcomes and can be held to account when those outcomes cause harm. It is not an after-the-fact investigation mechanism; it is a governance architecture designed before deployment. Organisations that don't map accountability before launching AI systems discover post-incident that nobody is formally responsible for anything — and the absence of formal responsibility is not a legal or ethical defence.

2. The Four Components of AI Accountability

Component	What It Requires	Failure Mode When Absent
Answerability	Responsible parties can explain decisions made in AI development and deployment	Post-incident investigations cannot reconstruct what happened or why
Enforceability	Mechanisms exist to impose consequences when AI causes harm — regulatory, legal, professional, or organisational	Harms occur without consequence, removing incentive for responsible development
Responsiveness	Affected individuals can reach accountable parties with complaints; defined response procedures exist	Users harmed by AI have no meaningful recourse; problems are not escalated
Remediation	When harm occurs, accountable parties take corrective action and prevent recurrence	The same harm recurs; no organisational learning happens

All four components must be present for accountability to function. Answerability without enforceability produces explanations without consequences. Enforceability without responsiveness makes accountability inaccessible to the people most affected. Responsiveness without remediation creates the appearance of accountability without the substance.

3. Governance Frameworks: Accountability Requirements

Every major AI governance framework dedicates substantial attention to accountability architecture. Understanding what each framework requires enables practitioners to map their organisation's obligations:

- NIST AI RMF — GV.RR (Roles and Responsibilities) requires defined ownership for AI risk decisions across the lifecycle: who approves deployment, who monitors production, who receives external concerns about AI harms. GV.OV (Oversight) requires that AI risk management processes are monitored and reviewed.
- ISO/IEC 42001:2023 Section 5 (Leadership) — requires top management to demonstrate commitment to AI governance and to assign accountability explicitly, not by implication.
- EU AI Act — defines two distinct legal roles: Provider (develops or places AI on the market) and Deployer (uses AI in practice), with separate obligations for each. This distinction prevents accountability from shifting between parties without formal assignment.
- GDPR Article 22 — provides a right to human review for significant automated decisions, creating a mechanism for attaching individual accountability to algorithmic decisions that would otherwise be fully automated.

EU AI Act provider vs. deployer distinction: The EU AI Act deliberately separates obligations between the organisation that builds AI and the organisation that uses it. Security practitioners should identify which role their organisation occupies for each AI system — the obligations, penalties, and documentation requirements differ substantially.

4. The AI Accountability Stack: Four Layers

Accountability does not reside in a single role or function — it operates simultaneously across multiple organisational layers. Effective governance maps each layer's accountability explicitly:

- Executive accountability — the C-suite and board bear strategic accountability for AI risk: which risks to accept, what governance resources to invest, and what ethical lines the organisation will not cross. These are governance decisions, not technical ones.
- Operational accountability — product managers, data scientists, and engineers bear accountability for specific implementation decisions: how the model was trained, what testing was done, what limitations were documented and disclosed.
- Governance accountability — ethics committees, legal teams, and compliance functions bear accountability for oversight: did they apply appropriate scrutiny? Did they receive the information they needed to exercise judgement?
- External accountability — regulators, auditors, and courts hold organisations accountable when harm occurs. This layer activates when internal accountability mechanisms have failed or are insufficient.

Mapping which layer bears which accountability for each AI system is a prerequisite for governance maturity. Without this mapping, incidents surface as disputes about whose fault it is rather than processes for remediation and learning.

5. Documentation as Accountability Infrastructure

Accountability requires evidence. The documentation chain for AI systems is the organisational memory that makes post-incident investigation, regulatory response, and organisational learning possible. Without it, accountability is a principle without a mechanism.

- Training data provenance — where did training data come from? What biases might it contain? Was it lawfully obtained? Was consent appropriate?
- Model evaluation records — what tests were run, on what populations, with what results? What performance gaps were identified and what was done about them?
- Risk assessment documentation — what harms were identified in pre-deployment risk assessment, what mitigations were applied, and what residual risks were formally accepted by whom?
- Deployment approval records — who authorised deployment and what conditions were attached? Were those conditions met?
- Incident records — when AI systems caused harm, what happened, what was done in response, and what changes were made to prevent recurrence?

This documentation should be maintained with the same rigour as information security audit logs. It is not supplementary record-keeping — it is the evidentiary foundation of accountability governance. Regulatory inquiries cannot be answered, legal proceedings cannot be navigated, and organisations cannot learn from failures without it.

6. Legal Accountability: An Evolving Landscape

The legal accountability landscape for AI harm has developed substantially since 2023 and continues to evolve. Practitioners operating in regulated environments must track these developments:

Framework	Key Accountability Provision	Maximum Penalty
EU AI Act (2024)	Provider and deployer obligations for high-risk AI; conformity assessments; human oversight requirements	€30M or 6% of global annual turnover
EU AI Liability Directive (proposed)	Civil liability for AI harms; burden of proof provisions	Civil damages to affected parties
US OMB AI Guidance	Federal agency accountability for AI governance and impact assessments	Administrative; agency-specific
CFPB / OCC Guidance (Financial AI)	Explainability and non-discrimination requirements for AI in credit decisions	Regulatory enforcement; civil money penalties
FDA (Medical AI)	Premarket approval and post-market monitoring for AI/ML-based medical devices	Product recall; civil and criminal penalties

The direction of travel is consistent: legal accountability for AI harm is increasing, regulatory frameworks are maturing, and organisations that built AI programs without accountability governance are accumulating exposure as these frameworks develop enforcement mechanisms.

7. Accountability Gaps as Security Vulnerabilities

Real-World Pattern — SyRI Case (Netherlands, 2020): The SyRI welfare fraud detection system operated with diffuse accountability — no individual or team had clear ownership of the system's fairness or outcomes. This diffusion meant that individual concerns about the system's discriminatory impact had no clear escalation path. The court found that the lack of transparency and accountability was itself part of the rights violation. Diffuse accountability is not merely a governance gap; it is a condition that enables sustained harm to continue without triggering the oversight mechanisms that would identify and correct it.

From a security perspective, accountability gaps create direct vulnerabilities. When ownership of AI systems is diffuse, oversight degrades, monitoring gaps appear, and incident response is slower because no one is certain who should lead the response. Security practitioners should treat accountability mapping as a component of AI system security architecture:

- Unmanaged AI systems — systems without a named owner — are analogous to unmanaged network devices: invisible to governance, opaque to investigation, and persistently exploitable.
- Insider threats who want to manipulate AI behaviour benefit from accountability gaps: when nobody has clear ownership of a model in production, unauthorised modifications may go undetected.
- Diffuse accountability creates slow incident response: when an AI security system misbehaves, the response time is directly affected by how quickly the accountable owner can be identified and engaged.

8. Automated Decisions and Accountability Challenges

Specific accountability challenges arise when AI decisions are automated without human review. When a human makes a discriminatory lending decision, that individual bears identifiable legal and professional accountability. When an algorithm makes the same decision at scale, accountability disperses:

- The developer didn't make individual decisions — they built a system that made patterns of decisions.
- The deployer may claim the algorithm was providing recommendations, not binding determinations.
- Regulators face practical difficulty assigning enforceable responsibility for aggregate effects of algorithmic decision-making.

Governance frameworks address this gap through two mechanisms. GDPR Article 22 establishes a right to human review for significant automated decisions, creating a point in the process where individual accountability attaches — a human must review and be responsible for acting on the AI recommendation. The EU AI Act's human oversight requirements for high-risk AI address the gap more systematically by requiring that humans can understand, monitor, and intervene in AI decision-making. Accountability for automated decisions requires either human review points where individual responsibility clearly attaches, or organisational accountability structures that formally accept system-level responsibility for aggregate algorithmic effects.

9. Practical Accountability Mechanisms

Accountability is not just a governance principle — it requires operational mechanisms that make it real. These are the practices that convert accountability policy into accountability function:

- Designated AI owners — named individuals formally responsible for each AI system in production, with defined authority and obligations. A system without a named owner is an unmanaged asset.
- Formal approval gates — defined checkpoints where responsible parties must sign off before AI advances through development stages. Each sign-off is a documented accountability attachment point.
- RACI matrices — documenting which role is Responsible, Accountable, Consulted, and Informed for each AI lifecycle activity. RACI maps convert implicit assumptions about accountability into explicit, auditable assignments.
- Incident accountability procedures — defining who leads AI incident response, investigation, and remediation before incidents occur. Discovering this in the middle of an incident creates delays and disputes.
- External reporting channels — mechanisms for individuals to report AI harms to an accountable contact, with defined response procedures and escalation paths.

Organisations with mature accountability mechanisms respond to AI incidents faster and extract more organisational learning from them. The mechanisms don't prevent AI harms, but they ensure that when harms occur there are clear, pre-established paths to response and remediation.

10. Professional Accountability for AI Practitioners

Individual practitioners bear professional accountability that exists independently of organisational structures. This is an important point: organisational accountability gaps do not eliminate individual professional responsibility.

- A data scientist who identifies bias in a model and does not escalate it has made a professional choice with accountability implications — the ACM Code of Ethics addresses algorithmic harm explicitly.
- An engineer who deploys a model without adequate testing under time pressure bears accountability for the foreseeable consequences of inadequate testing.
- A security practitioner who discovers that an AI system is being used beyond its validated scope and does not report it has failed a professional obligation.
- IEEE ethics guidelines for autonomous and intelligent systems articulate the professional norms that define negligent AI practice.

These professional codes do not have the enforcement mechanisms of medical or legal regulation, but they define the standards against which professional judgement will be evaluated in post-incident review, litigation, and regulatory inquiry. Practitioners who document their escalation attempts when organisational culture suppresses concerns are protecting themselves professionally while meeting their ethical obligations.

11. Whistleblower Protections and Internal Escalation

Accountability governance requires that people who identify problems can raise them without career risk. Without this protection, the individuals best positioned to identify AI failures early — the teams closest to the systems — are incentivised to remain silent.

- Whistleblower protections apply in various jurisdictions to employees who report violations of law — including AI-related violations of GDPR, anti-discrimination law, or sector-specific regulation.
- Internal escalation paths — ethics hotlines, AI ethics committee reporting channels, ombudsperson offices — provide mechanisms for raising concerns below the threshold of formal regulatory reporting.
- Organisations where escalating ethics concerns is career-limiting are systematically degrading their own accountability mechanisms: they are eliminating the organisational immune system that would otherwise identify failures early.

The connection between whistleblower protection and accountability governance is direct: accountability requires information to flow to the people and functions responsible for acting on it. Culture that suppresses that information flow makes accountability structurally impossible, regardless of how well-designed the formal accountability architecture is.

Key Terms Glossary

Term	Definition
Accountability (AI)	Governance requirement that identifiable parties bear responsibility for AI outcomes and can be held to account when those outcomes cause harm, encompassing answerability, enforceability, responsiveness, and remediation.
Answerability	The ability of responsible parties to explain the decisions made in AI development and deployment, forming the evidentiary basis for accountability.
GV.RR (NIST AI RMF)	Roles and Responsibilities function requiring that policies, processes, and practices across the AI lifecycle have defined ownership.
Provider (EU AI Act)	The organisation that develops or places an AI system on the market; bears the primary compliance and conformity obligations under the EU AI Act.
Deployer (EU AI Act)	The organisation that uses an AI system in practice; bears operational obligations under the EU AI Act, distinct from those of the provider.
RACI Matrix	Accountability tool mapping Responsible, Accountable, Consulted, and Informed roles for each AI lifecycle activity, converting implicit accountability assumptions into explicit assignments.
SyRI	Dutch government AI system for welfare fraud detection, shut down by court order in 2020 for privacy and human rights violations — a landmark AI accountability case.
GDPR Article 22	Provision establishing a right to human review for significant automated decisions, creating a mechanism for attaching individual accountability to algorithmic decision-making.
AI Liability Directive (EU, proposed)	Proposed EU framework establishing civil liability for AI harms, with burden of proof provisions that shift the evidential challenge to AI providers and deployers.
Designated AI owner	Named individual formally responsible for a specific AI system in production, with defined authority and obligation — the

Quick Revision Checklist

- Accountability requires four components: answerability, enforceability, responsiveness, and remediation — all four must be present for the mechanism to function.
- NIST AI RMF GV.RR requires defined ownership for AI risk decisions; ISO 42001 Section 5 requires top management to explicitly assign accountability.
- The EU AI Act distinguishes provider and deployer roles with separate obligations — identify which role your organisation occupies for each AI system.
- GDPR Article 22 creates a right to human review for significant automated decisions, attaching individual accountability to algorithmic outcomes.
- AI documentation (training data provenance, evaluation records, risk assessments, deployment approvals, incident records) is the evidentiary foundation of accountability.
- EU AI Act penalties reach €30M or 6% of global annual turnover for the most serious violations — legal accountability is financially material.
- Diffuse accountability is a security vulnerability: unmanaged AI systems (no named owner) are invisible to governance and slower to incident response.
- Practitioners bear professional accountability independently of organisational structures — the ACM Code of Ethics and IEEE guidelines define the professional standard.
- Whistleblower protections and internal escalation paths are accountability infrastructure — culture that suppresses escalation makes accountability structurally impossible.
- Designated AI owners, formal approval gates, RACI matrices, and incident accountability procedures convert accountability policy into operational accountability function.

10 Exam-Based MCQs — Accountability

Test yourself after reading. These questions are written in exam style — scenario-heavy, with plausible distractors. Read every explanation, including for questions you answered correctly.

Q1. *Scenario:* A financial services company is deploying an AI credit scoring system. The CISO is working with the compliance team to implement NIST AI RMF governance for the deployment. The team must ensure that accountability is properly established before go-live. Which action BEST establishes accountability for AI deployment in accordance with NIST AI RMF GV.RR?

- A) Documenting the AI model's technical architecture in the system design specification
- B) Assigning named individuals with defined authority as accountable owners for each AI system and lifecycle decision point
- C) Conducting a post-deployment audit of AI outputs for bias within 90 days of launch
- D) Requiring all AI vendors to provide indemnification for AI-related harms

Answer: B **Explanation:** B is correct because GV.RR (Roles and Responsibilities) specifically requires that policies, processes, and practices across the AI lifecycle have defined ownership — named accountable individuals, not just roles or teams. A is wrong because technical architecture documentation is a design

artifact, not an accountability assignment. C is wrong because post-deployment auditing is an accountability mechanism but assigning the auditor and acting on findings still requires the named accountability that GV.RR requires. D is wrong because vendor indemnification is a contractual risk transfer mechanism, not an internal accountability assignment — the organisation still needs to assign accountability for its own deployment decisions.

Q2. Scenario: A retail bank's AI loan approval system has been found to systematically deny loans at higher rates to applicants from a specific postcode area, correlating with a minority demographic. The error has affected 3,200 applicants over 18 months. The bank's CISO has escalated to the AI governance committee. Which response BEST demonstrates appropriate accountability governance for this incident?

- A) Issue a statement that the AI system was operating as designed and the bias is a training data issue outside the company's control
- B) Activate the incident accountability procedure, identify the accountable AI owner, conduct a root cause investigation, remediate affected individuals, and implement controls to prevent recurrence
- C) Disable the AI system immediately and revert to manual processing for all decisions
- D) Refer the matter to the AI vendor for investigation and await their findings before taking internal action

Answer: B Explanation: B is correct because effective accountability governance requires all four accountability components: answerability (root cause investigation), enforceability (taking responsibility rather than deflecting), responsiveness (remediating affected individuals), and remediation (controls to prevent recurrence). A is wrong because deflecting accountability to training data — while potentially partially accurate — does not address the organisation's deployment responsibility or the affected individuals' need for redress. C is wrong because immediately disabling the system, while potentially appropriate, is not itself an accountability response — it must be accompanied by the investigation and remediation steps. D is wrong because delegating investigation to the vendor delays the accountability response and doesn't address the deployer's independent obligations to affected individuals.

Q3. Scenario: A European bank has licensed an AI credit decisioning system from a US-based AI provider. An EU regulator has opened an investigation into potential violations of the EU AI Act related to the system's use of protected characteristics. The bank's legal team is determining which obligations apply to the bank versus the AI vendor. How does the EU AI Act's distinction between provider and deployer MOST directly affect the bank's accountability obligations?

- A) The bank has no obligations because it is a deployer, not the AI system's provider
- B) As a deployer, the bank has specific operational obligations — including human oversight, use within intended purpose, and monitoring — that are separate from the provider's obligations
- C) The bank and the AI vendor share equal accountability, so obligations can be divided between them by contract
- D) The EU AI Act only applies to providers; deployers are governed exclusively by national law

Answer: B Explanation: B is correct because the EU AI Act explicitly assigns separate obligations to deployers of high-risk AI, including ensuring human oversight, using the system within its intended

purpose, monitoring outputs, and informing affected individuals of AI-based decisions. Deployers cannot transfer these obligations to providers. A is wrong because deployer status creates obligations, not immunity — the Act is specifically designed to prevent accountability gaps that arise when deployers claim they are just using someone else's system. C is wrong because while contracts can manage commercial risk allocation between provider and deployer, they cannot transfer statutory obligations imposed by the AI Act. D is wrong because the EU AI Act explicitly addresses both providers and deployers in its compliance framework.

Q4. Which of the following BEST describes what a RACI matrix provides in AI accountability governance?

- A) A technical framework for testing AI model accuracy across demographic groups
- B) A documentation of which roles are Responsible, Accountable, Consulted, and Informed for each AI lifecycle activity
- C) A regulatory reporting template for AI incident disclosure
- D) A risk scoring methodology for AI deployment decisions

Answer: B **Explanation:** B is correct because a RACI matrix is an accountability assignment tool that converts implicit assumptions about who is responsible for what into explicit, documented, and auditable assignments. It directly addresses the diffuse accountability problem by making ownership of each AI lifecycle activity unambiguous. A is wrong because RACI is a governance tool, not a testing framework. C is wrong because RACI is an internal accountability document, not a regulatory reporting template. D is wrong because RACI assigns roles and responsibilities rather than scoring deployment risks.

Q5. A security practitioner discovers that an AI threat detection tool deployed in production is being used in a broader set of contexts than it was validated for. What is the practitioner's MOST appropriate action?

- A) Continue using the tool, as expanding AI use beyond validated scope is a normal part of operational AI use
- B) Immediately disable the tool pending a formal risk assessment
- C) Document and report the out-of-scope use through the organisation's AI governance escalation channel
- D) Contact the AI vendor to request retroactive validation of the expanded scope

Answer: C **Explanation:** C is correct because the practitioner bears professional accountability for identifying and escalating the out-of-scope use — documenting and reporting through governance channels fulfils this obligation and activates the accountability mechanisms responsible for reviewing and deciding on the expanded use. A is wrong because using AI beyond its validated scope is a governance failure, not a normal operational choice; the practitioner has a professional obligation to report rather than normalise it. B is wrong because immediately disabling the tool may be disproportionate and is an operational decision that exceeds the practitioner's individual authority; escalating through governance channels is the appropriate first step. D is wrong because retroactive validation does not address the governance failure that occurred during the period of out-of-scope use, and the vendor cannot be the primary accountability mechanism for the deploying organisation's decisions.

Q6. Under GDPR Article 22, what specific accountability mechanism does the right to human review for automated decisions create?

- A) It requires organisations to publish the technical parameters of any AI system making automated decisions
- B) It creates a point in the automated decision process where a human must review and bear identifiable accountability for the AI-informed decision
- C) It prohibits fully automated decision-making for any decision affecting EU residents
- D) It requires AI providers to register automated decision-making systems with national data protection authorities

Answer: B **Explanation:** B is correct because Article 22 addresses the accountability gap in fully automated decisions by requiring human review for decisions that significantly affect individuals — this creates a point where individual human accountability attaches to what would otherwise be a diffuse algorithmic process. A is wrong because Article 22 is about decision-making accountability, not technical parameter disclosure (that is addressed by transparency and explainability requirements elsewhere in GDPR). C is wrong because Article 22 does not prohibit fully automated decision-making; it creates a right for individuals to request human review, with exceptions for contractual necessity and explicit consent. D is wrong because Article 22 creates individual rights, not a registration requirement for providers.

Q7. Which of the following accountability documentation items is MOST critical for enabling post-incident investigation of an AI system failure?

- A) The AI system's marketing materials and user-facing product description
- B) Training data provenance, model evaluation records, deployment approval records, and incident records
- C) The development team's sprint retrospective notes
- D) The AI vendor's standard terms of service agreement

Answer: B **Explanation:** B is correct because post-incident investigation requires reconstructing what data was used, what testing was done, who approved deployment and under what conditions, and what has been observed in production — these are exactly the items in the accountability documentation chain. A is wrong because marketing materials describe external positioning, not the development and deployment decisions that post-incident investigation needs to reconstruct. C is wrong because sprint retrospective notes are informal process documentation that rarely capture the specific accountability-relevant decisions about training data, testing, and deployment approval. D is wrong because vendor terms of service address commercial obligations, not the deploying organisation's internal accountability evidence chain.

Q8. Which of the following BEST describes why diffuse AI accountability is a security vulnerability?

- A) Diffuse accountability reduces AI model accuracy, creating security gaps in AI-driven threat detection

- B) When no individual or team has clear AI system ownership, oversight degrades, monitoring gaps emerge, and incident response is slower — creating conditions that insiders and attackers can exploit
- C) Diffuse accountability leads to excessive documentation, which increases the attack surface for data exfiltration
- D) Diffuse accountability reduces AI model performance, creating higher false positive rates in security tools

Answer: B **Explanation:** B is correct because the security risk of diffuse accountability is precisely that it creates management blind spots: unmanaged AI systems are invisible to governance, monitoring is degraded when ownership is unclear, and incident response delays give attackers additional time. The SyRI case demonstrated that diffuse accountability allowed a problematic system to continue operating without triggering the oversight mechanisms that should have identified the problem earlier. A is wrong because accountability diffusion is a governance failure, not a technical factor affecting model accuracy. C is wrong because documentation quantity does not in itself increase data exfiltration attack surface — that is a data management security question unrelated to accountability governance. D is wrong because accountability governance does not directly affect model inference performance.

Q9. An AI engineer discovers that a credit scoring model being prepared for deployment has a bias problem that will affect minority applicants. The engineer raises the concern with their manager, who instructs them to proceed with deployment on schedule. What does professional accountability require the engineer to do?

- A) Follow the manager's instruction — operational accountability rests with management, not individual engineers
- B) Immediately report the problem publicly to regulators without further internal escalation
- C) Escalate the concern through available internal governance channels (ethics committee, legal, compliance) and document the escalation attempt
- D) Refuse all further participation in the project and resign

Answer: C **Explanation:** C is correct because professional accountability (reflected in the ACM Code of Ethics and IEEE ethics guidelines) requires practitioners to escalate identified harms through available channels. Internal escalation to governance functions (ethics committee, legal, compliance) is both the professional obligation and the appropriate procedural step before external reporting. Documentation of the escalation attempt protects the practitioner professionally if the harm occurs despite escalation. A is wrong because professional accountability is individual and exists independently of managerial direction — following an instruction that results in foreseeable harm does not extinguish individual accountability. B is wrong because external regulatory reporting is appropriate when internal channels have been exhausted or are unavailable, not as an initial step; the practitioner should exhaust internal options first. D is wrong because resignation does not fulfil the professional obligation to prevent foreseeable harm — it merely removes the practitioner from involvement.

Q10. Which statement MOST accurately describes the trend in legal accountability for AI harm?

- A) Legal accountability for AI harm remains primarily a matter of contractual negotiation between vendors and deployers
- B) Legal accountability frameworks for AI are maturing globally, with increasing penalties, defined roles, and civil liability mechanisms — creating material financial exposure for organisations without accountability governance
- C) Legal accountability for AI harm applies only to AI providers, not to organisations that deploy commercially available AI
- D) Legal accountability for AI harm is limited to jurisdictions that have enacted specific AI legislation

Answer: B **Explanation:** B is correct because the trend across the EU AI Act, proposed EU AI Liability Directive, US sector-specific guidance, and national AI laws is consistent: legal accountability is increasing, enforcement mechanisms are developing, penalties are significant (EU AI Act up to 6% of global annual turnover), and organisations that built AI programs without accountability governance are accumulating exposure. A is wrong because while contracts manage commercial risk allocation, statutory obligations imposed by AI regulations cannot be contracted away between private parties. C is wrong because the EU AI Act explicitly assigns obligations to deployers, and sector-specific regulation (CFPB, FDA, EEOC) creates deployer accountability independently of provider status. D is wrong because sector-specific regulation creates accountability in many jurisdictions regardless of whether specific AI legislation exists — anti-discrimination law, consumer protection law, and data protection regulation all create accountability for AI-driven harms.